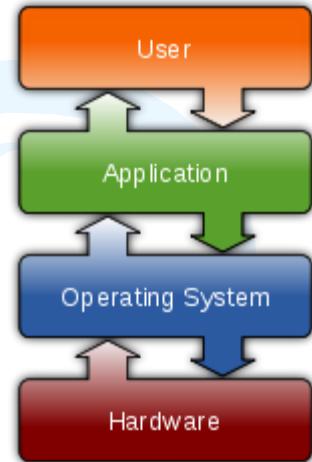


ENCE360

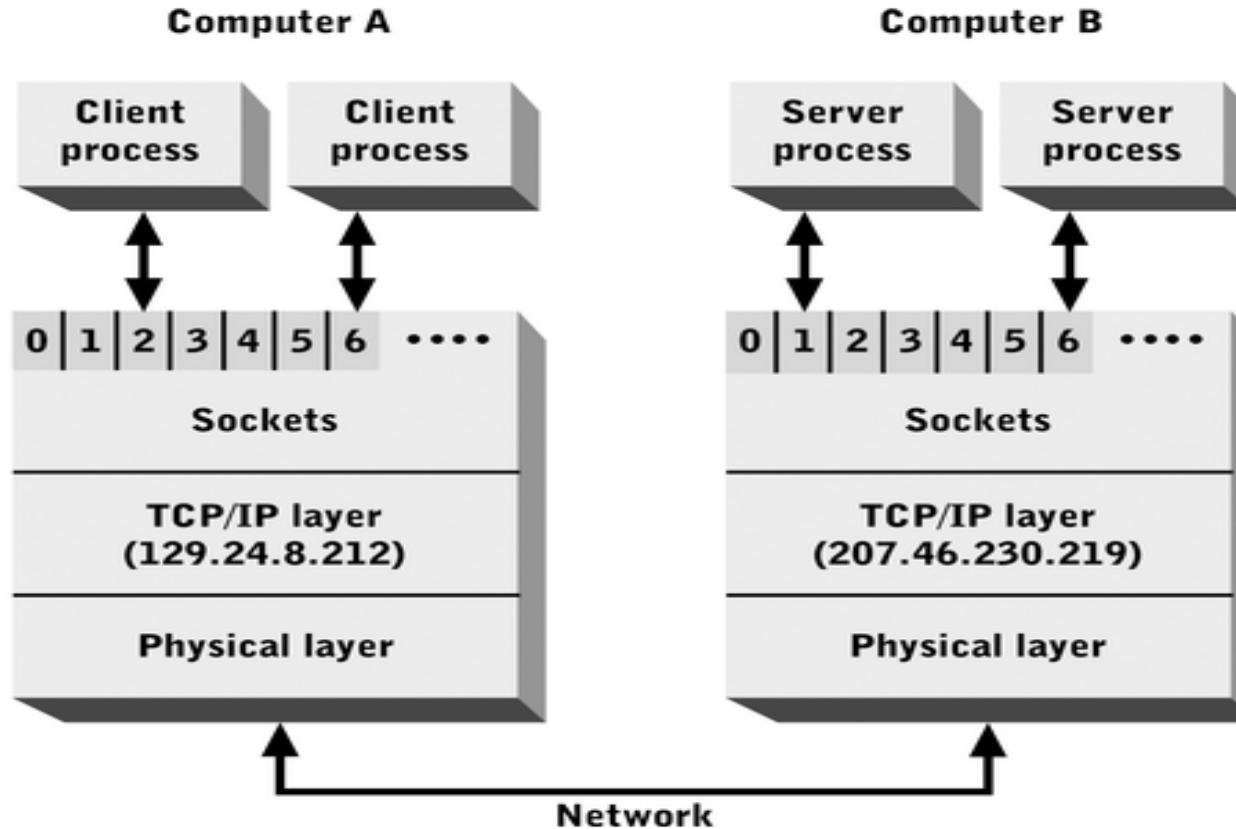
Operating Systems



**Interprocess
Communication**

Sockets

Interprocess Communication Using Sockets



(TCP=Transport Control Protocol, IP=Internet Protocol)



Sockets

Socket API

- introduced in
BSD4.1 UNIX,
1981

Two types of
sockets

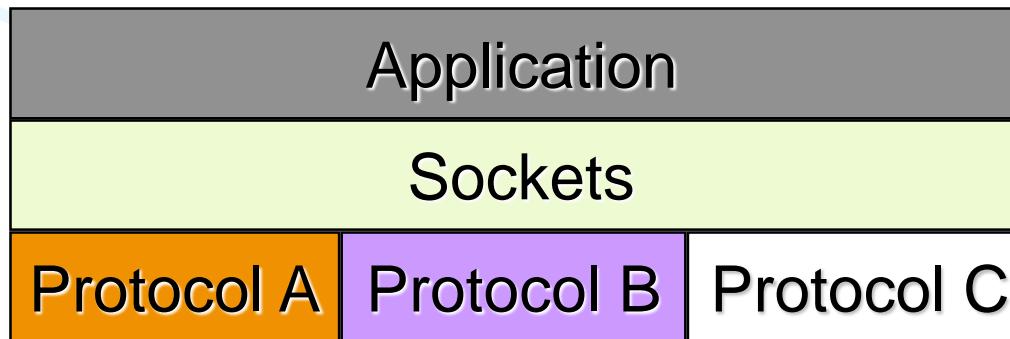
- connection-oriented
(stream)
- Connectionless
(datagram)

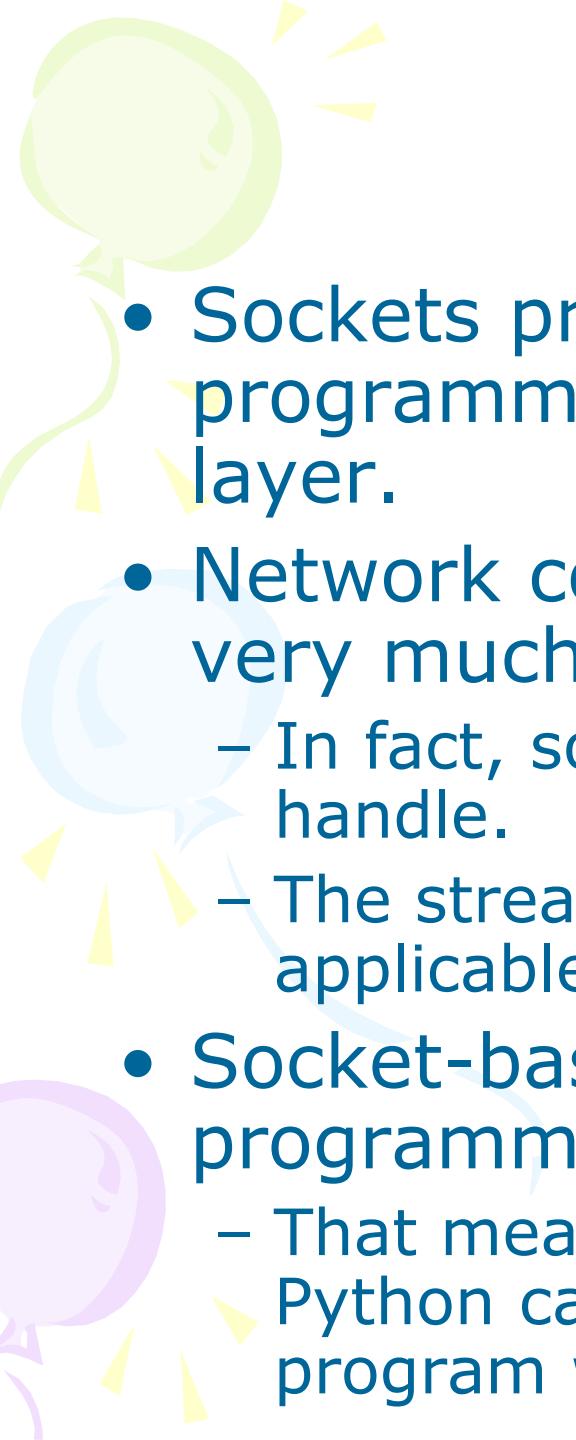
socket —

an interface
into which one
application process can
both send and
receive messages to/from
another (remote or
local) application process

What is a socket?

- API: network programming interface
- Socket: abstraction through which an application may send/receive data.
- Allows an application to “plug-in” a network and communicates with other applications in the same network.
- Sockets: support different underlying protocol families.





Sockets

- Sockets provide an interface for programming networks at the transport layer.
- Network communication using Sockets is very much similar to performing file I/O
 - In fact, socket handle is treated like file handle.
 - The streams used in file I/O operation are also applicable to socket-based I/O
- Socket-based communication is programming language independent.
 - That means, a socket program written in Python can also communicate with a socket program written in C or Java.

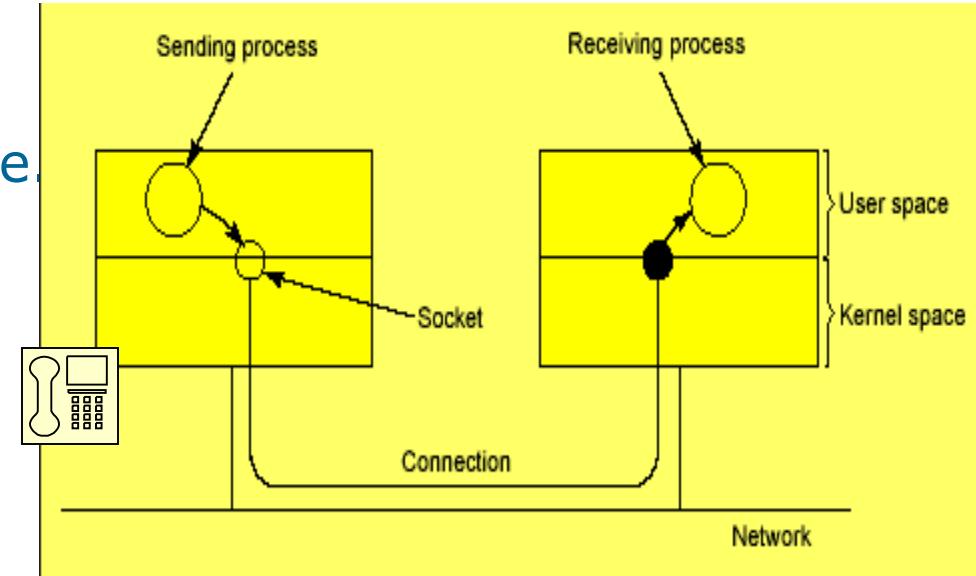
Socket

A socket is one endpoint of a two-way communication link between two programs running on the network.

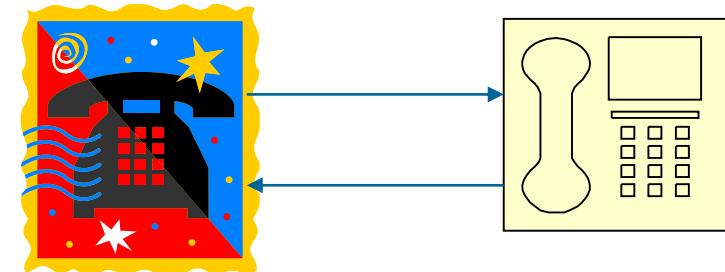
A socket is bound to a port number so that the network can identify the application that data is destined to be sent.

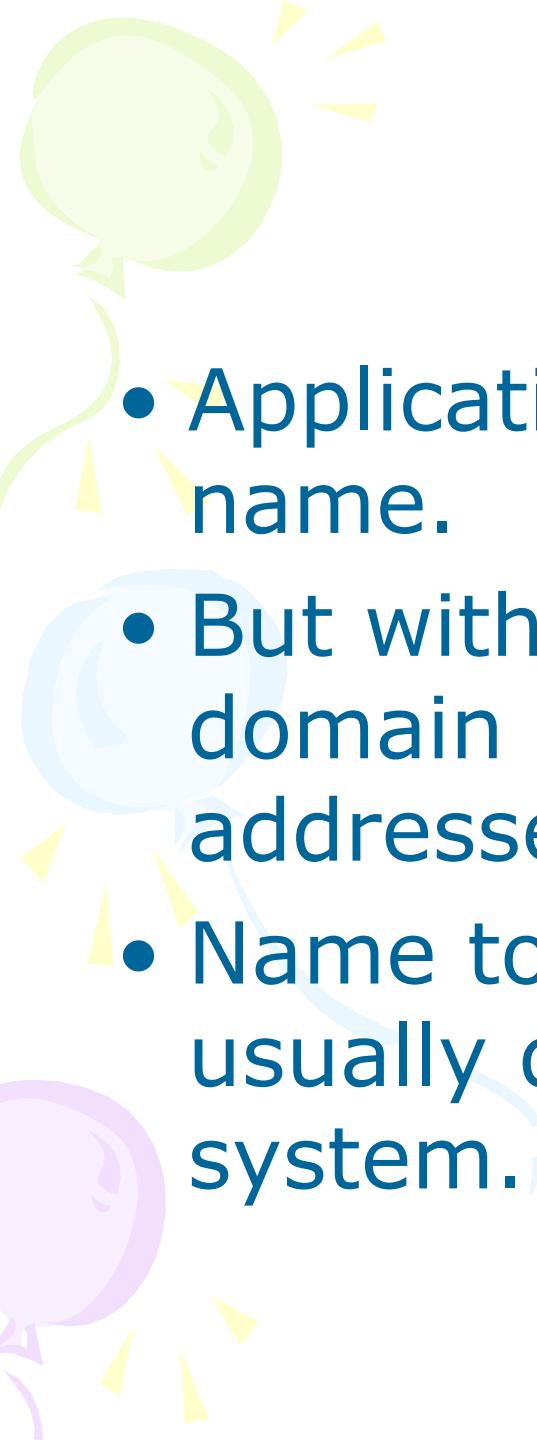
- Sockets connect unrelated processes that may be on different computers.
 - Once established, the socket connection may be read and written like a pipe.
- There are different domains of communication:
 - **UNIX/LINUX domain**
 - Sockets have actual file names. They can only be used with processes that reside on the same host.
 - Type is **AF_UNIX**
 - **Internet domain**
 - Allows unrelated processes different hosts to communicate via the Internet
 - Type is **AF_INET**

Socket basics



Analogy:





Socket Names

- Applications refer to sockets by name.
- But within the communication domain sockets are referred by addresses.
- Name to address translation is usually done outside the operating system.

A Socket-eye view of the Internet



atlas.cs.uga.edu
(128.192.251.4)



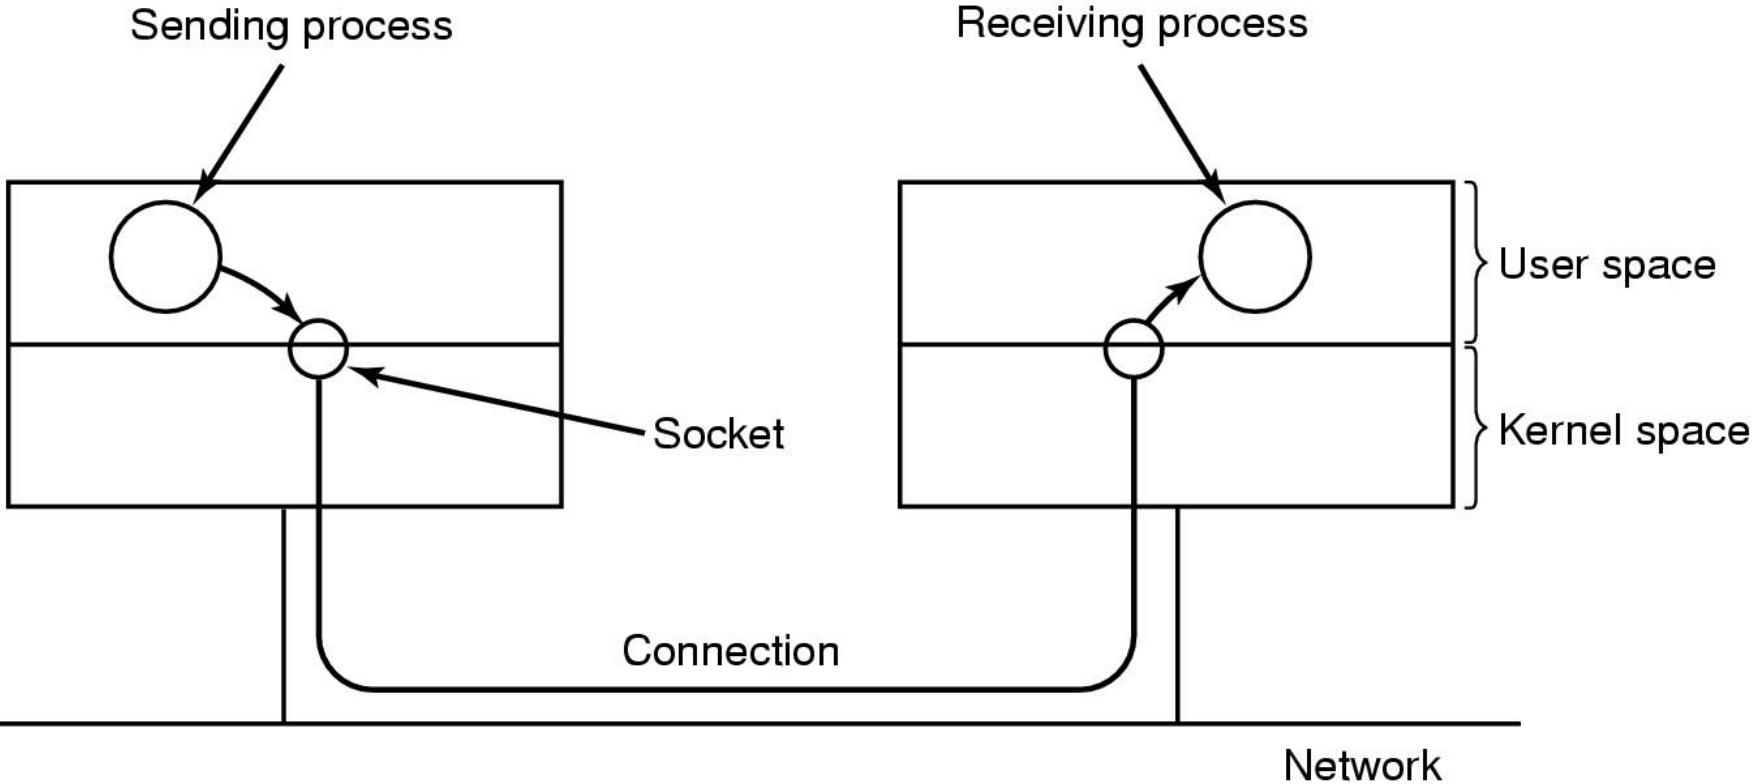
www.google.com
(216.239.35.100)



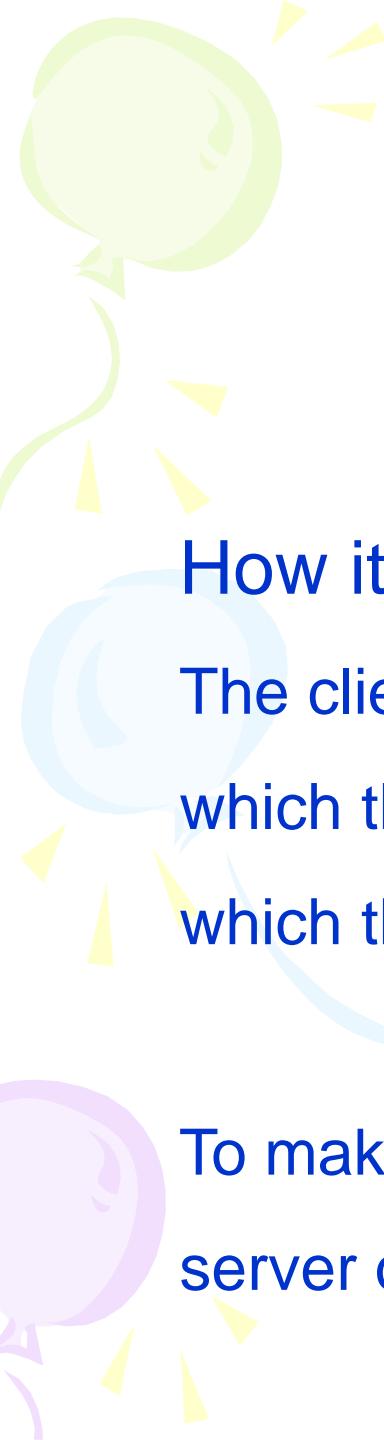
church.cse.ogi.edu
(129.95.50.2, 129.95.40.2)

- Each host machine has an IP address

Networking



Use of **s**ockets for networking



Socket

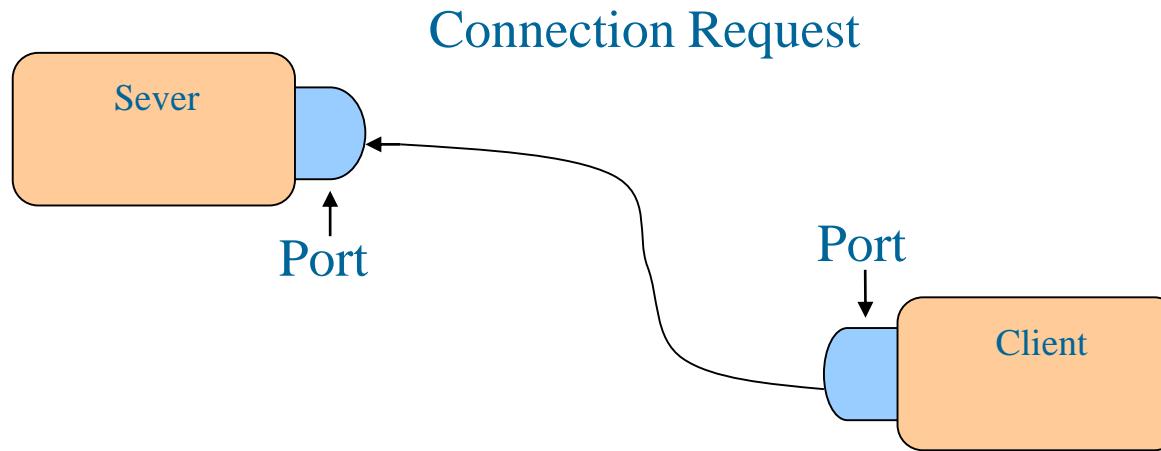
How it works?

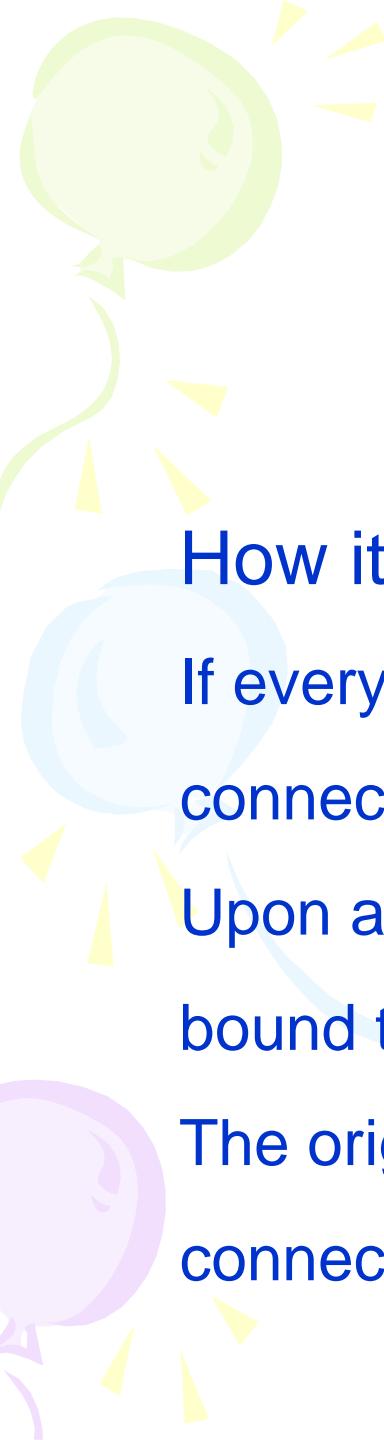
The client must know the hostname of the machine on which the server is running and the port number to which the server is connected.

To make a connection, the client sends a request to the server on the server's machine and port.

Socket

How it works?





Socket

How it works?

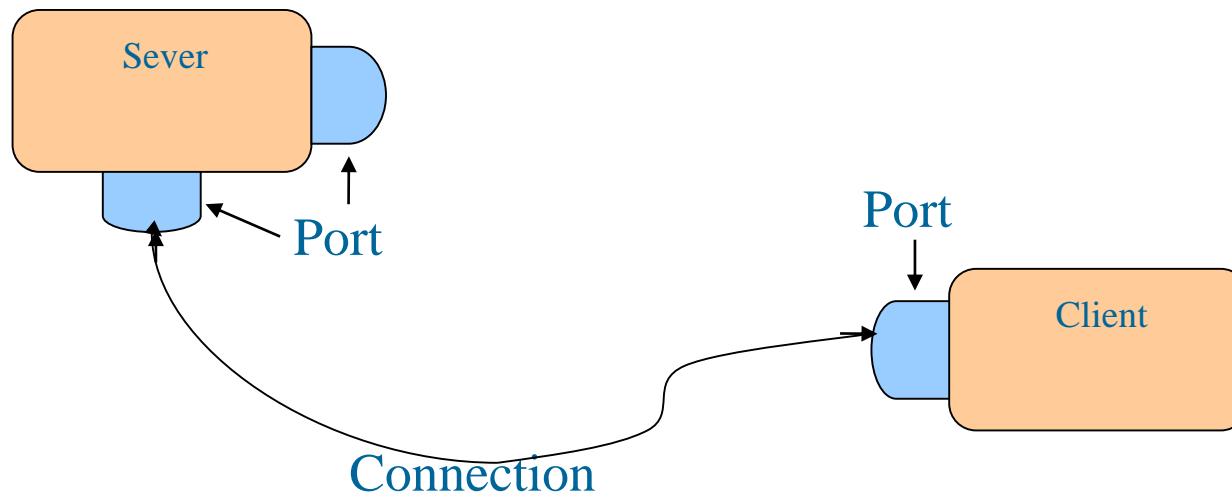
If everything goes well, the server accepts the connection.

Upon acceptance, the server creates a new socket bound to a different port.

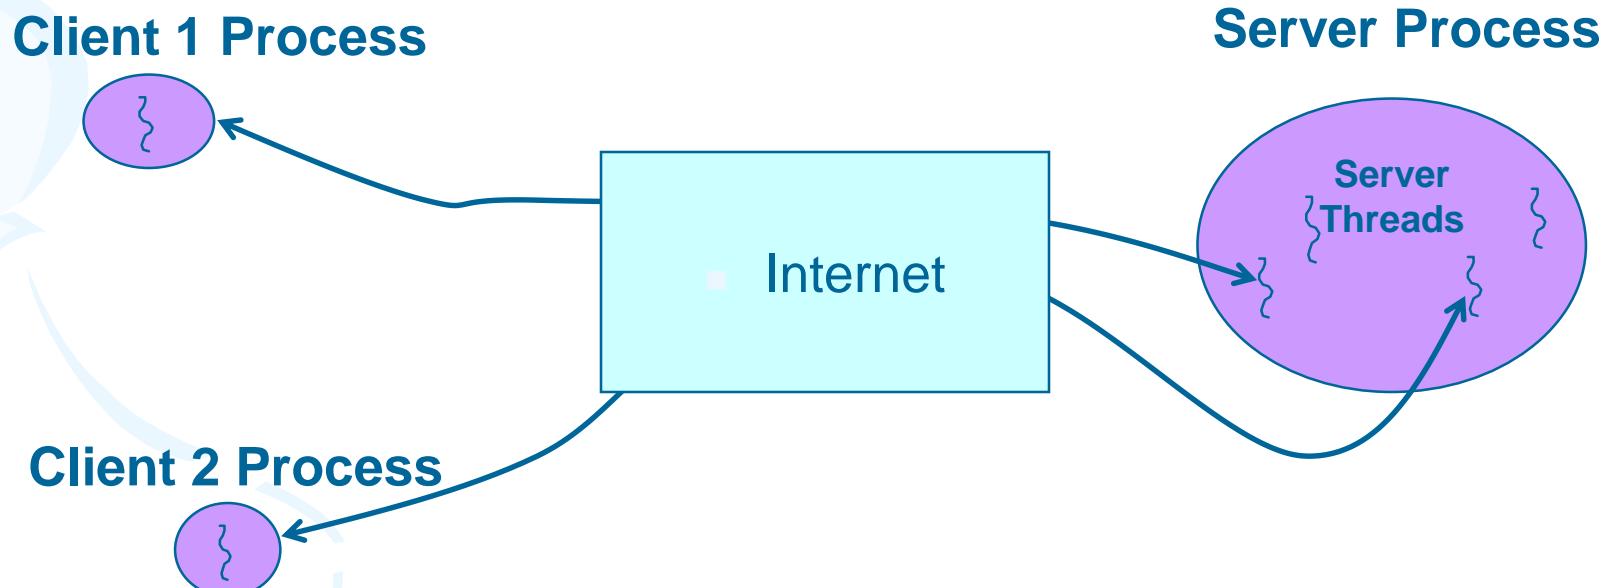
The original server port will be ready to listen for other connection request to the server.

Socket

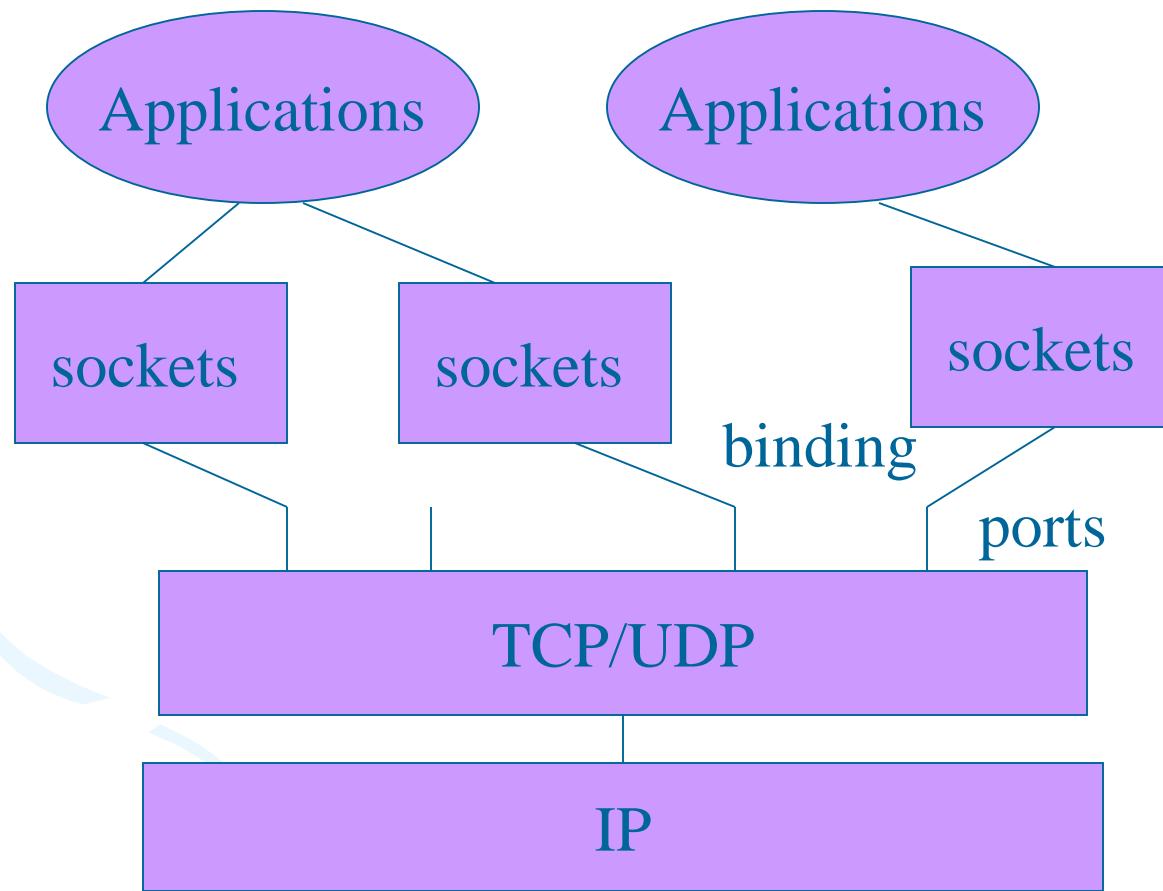
How it works?



Multithreaded Server: For Serving Multiple Clients Concurrently



Sockets, protocols, ports



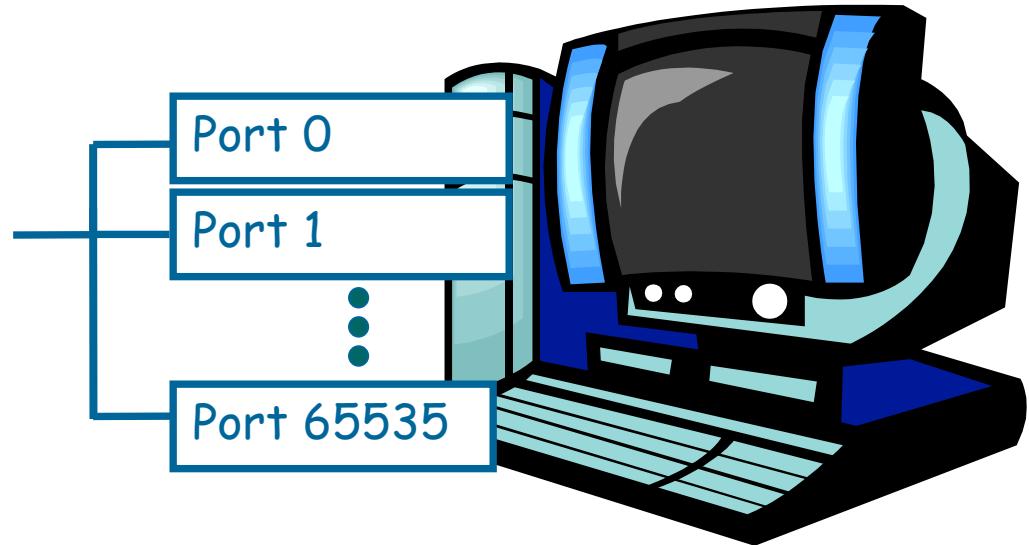
(**TCP**=Transport Control Protocol, **UDP**=User Datagram Protocol **IP**=Internet Protocol)

Ports

- Each host has 65,536 ports
- Some ports are *reserved for specific apps*
 - 20,21: FTP
 - 23: Telnet
 - 80: HTTP

about 1023 ports are reserved

User level process/services generally use port number value ≥ 1024



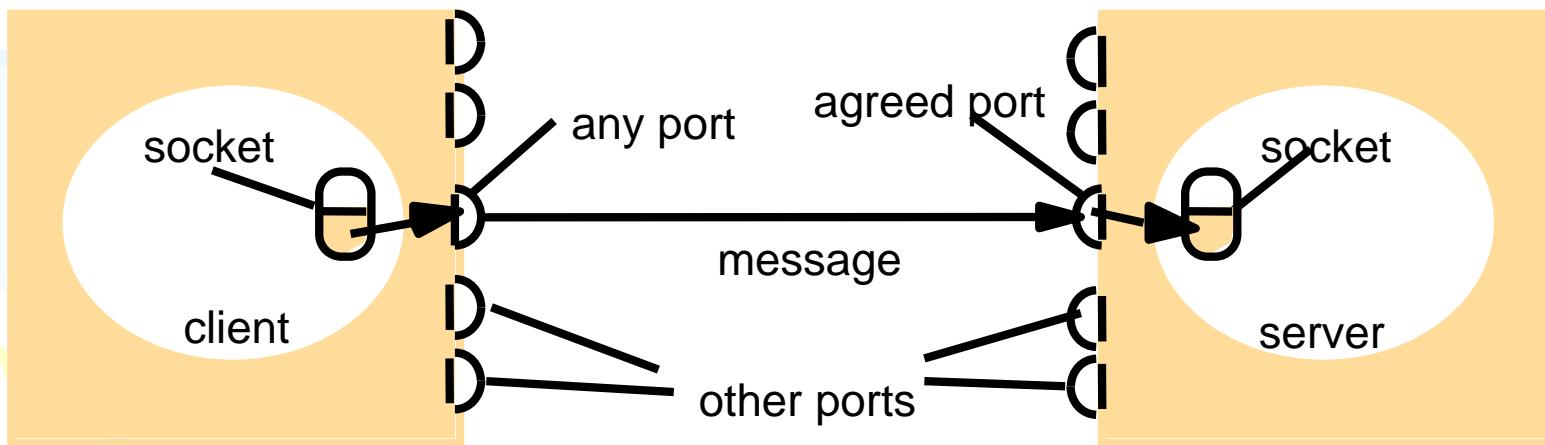
- A socket provides an interface to send data to/from the network through a port



Addresses, Ports and Sockets

- Like apartments and mailboxes
 - You are the application
 - Your apartment building address is the address
 - Your mailbox is the port
 - The post-office is the network

Sockets and ports

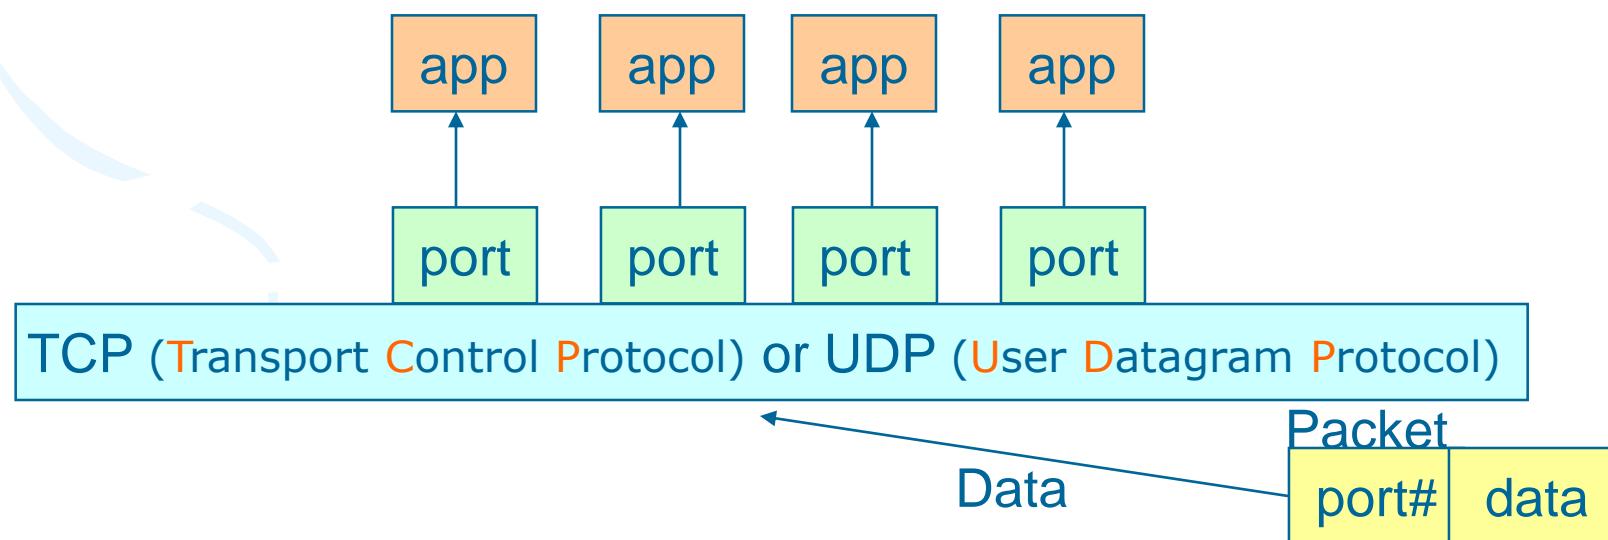
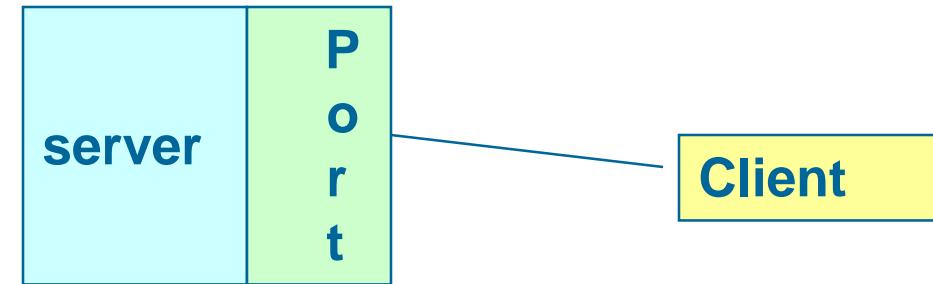


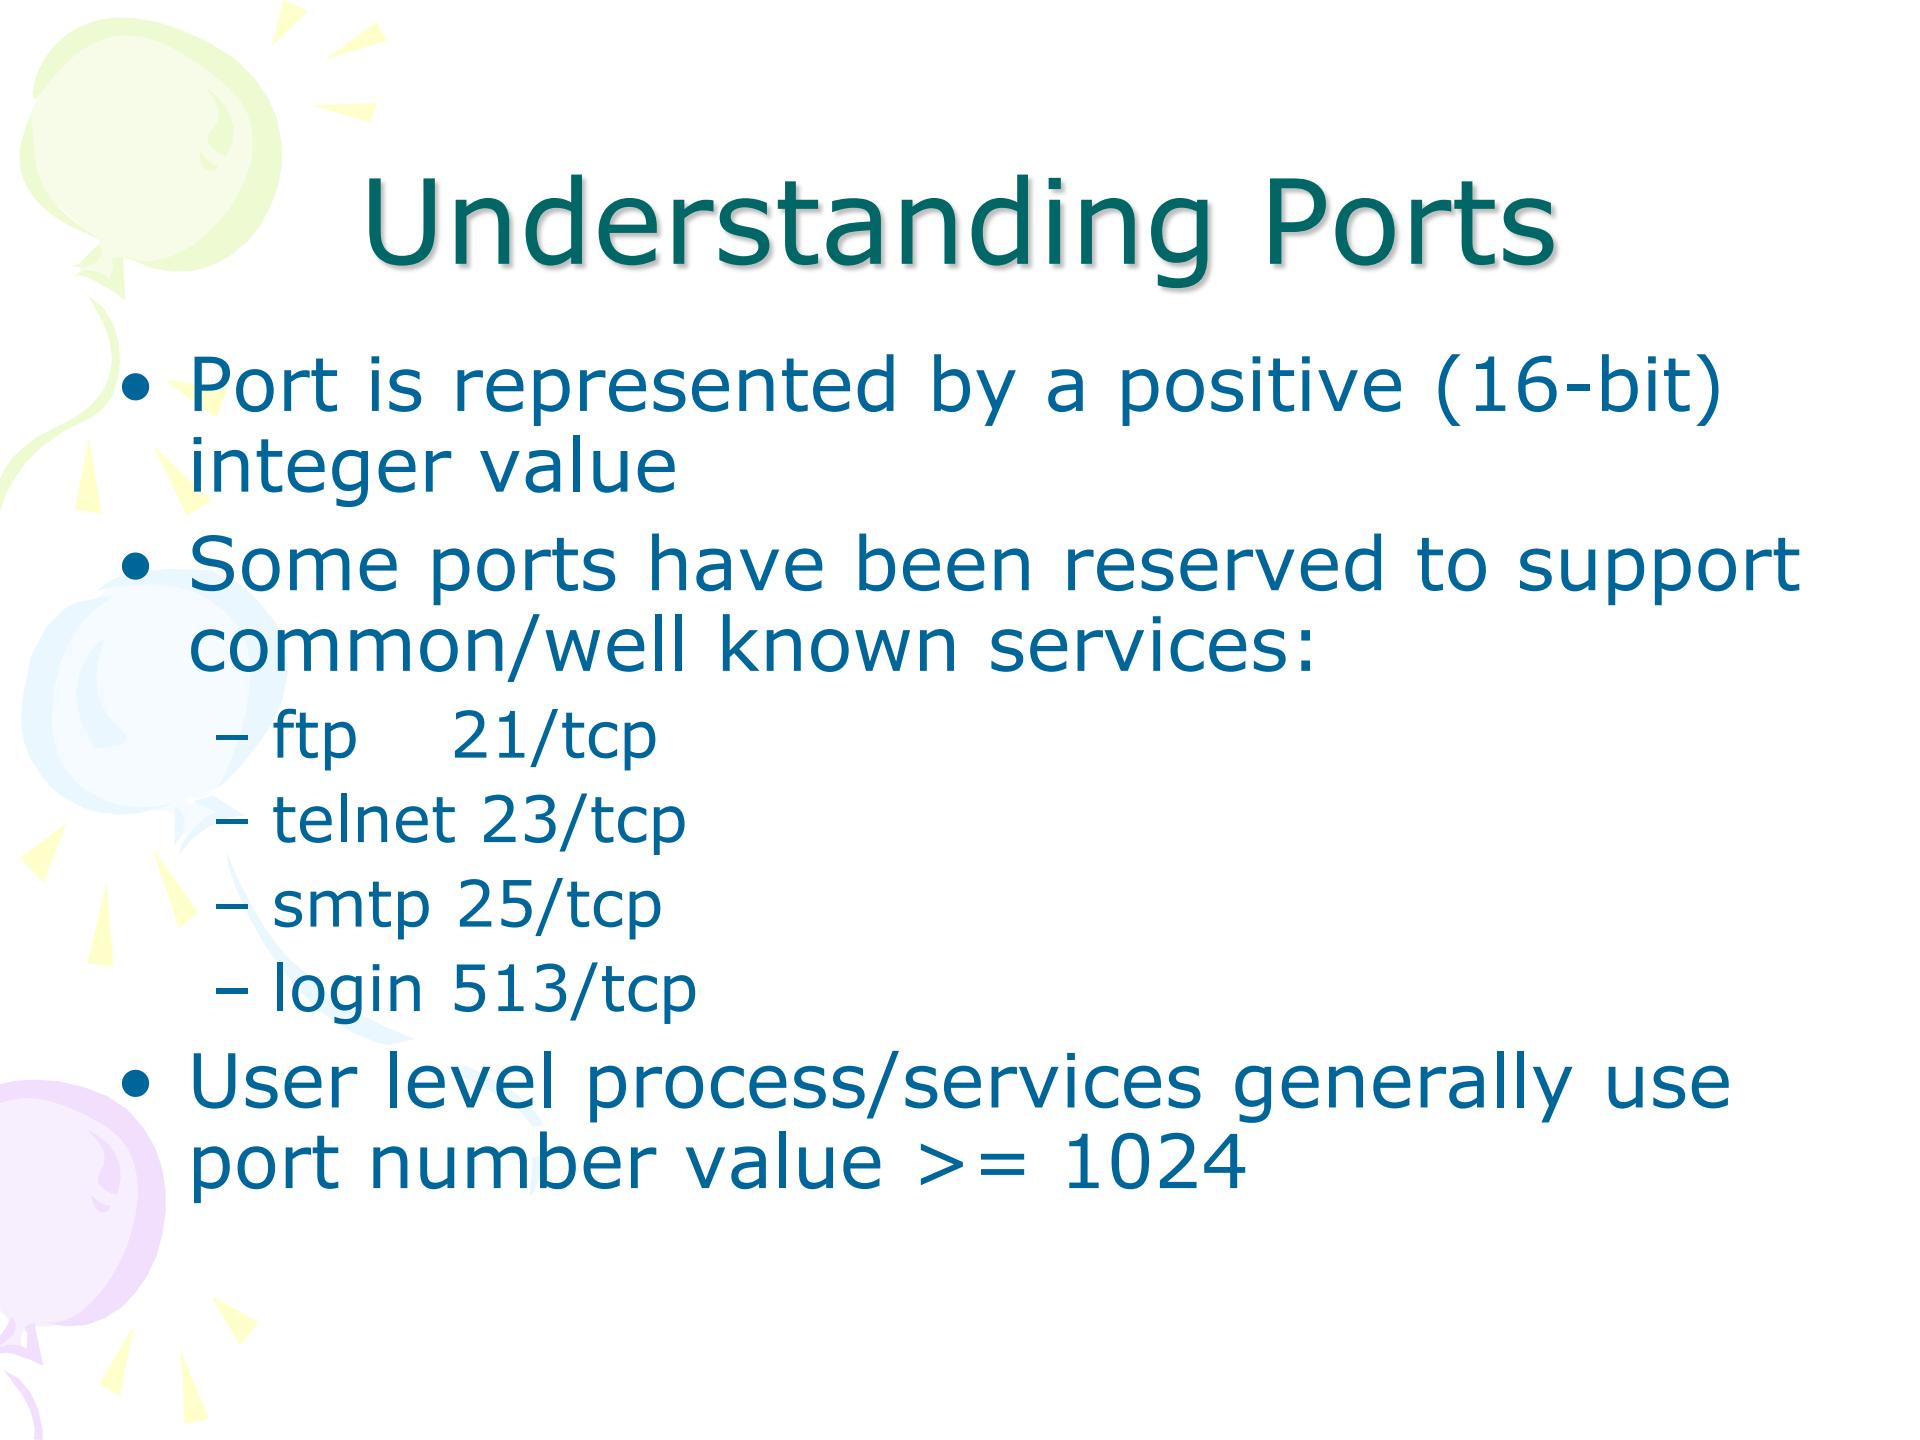
Internet address = 138.37.94.248

Internet address = 138.37.88.249

Understanding Ports

- Sockets use *ports* to map incoming data to a particular *process* running on a computer.





Understanding Ports

- Port is represented by a positive (16-bit) integer value
- Some ports have been reserved to support common/well known services:
 - ftp 21/tcp
 - telnet 23/tcp
 - smtp 25/tcp
 - login 513/tcp
- User level process/services generally use port number value ≥ 1024



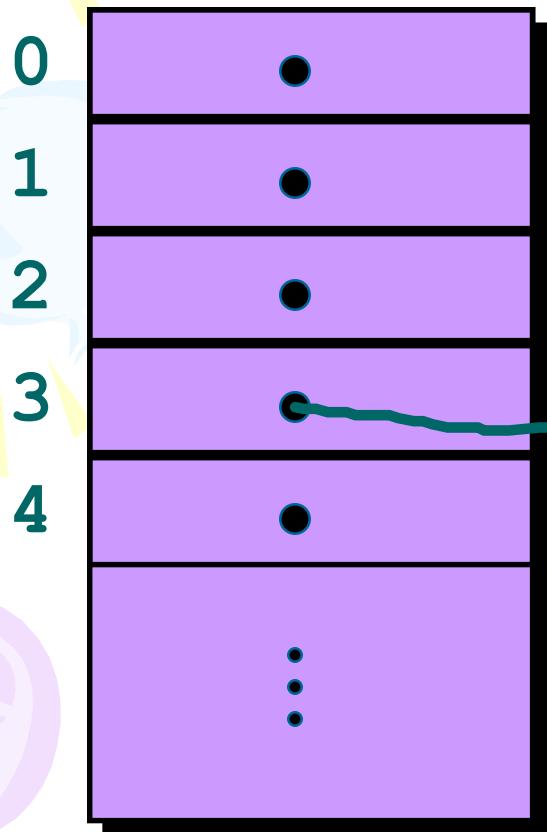
Standard Web Protocols & Services

All resources are identified by a unique **Uniform Resource Locator** (URL). The URL has four parts: protocol, host, port and resource.

- **Protocol** – an optional header specifying the resource access protocol.
- **Host** – the IP number or registered name of an Internet host computer or device.
- **Port** – an optional port number that specifies the **socket**.
- **Resource** – the complete path name of a resource on the host.

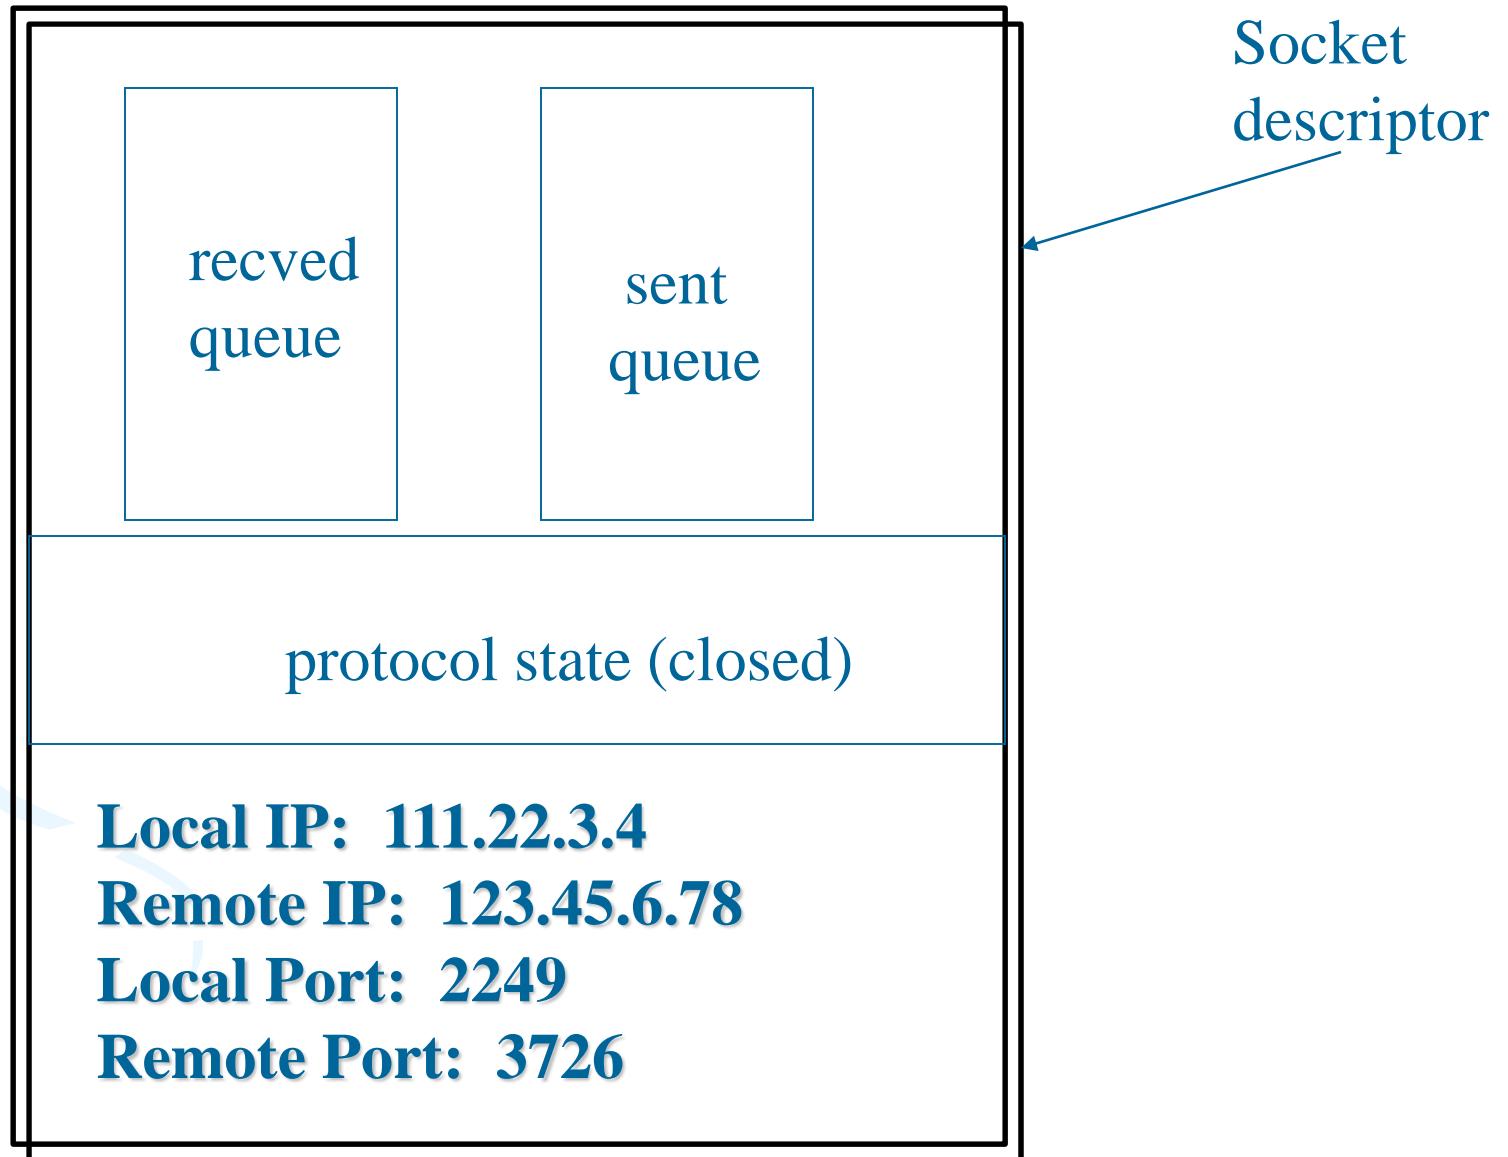
Socket Descriptor Data Structure

Descriptor Table



Family: PF_INET
Service: SOCK_STREAM
Local IP: 111.22.3.4
Remote IP: 123.45.6.78
Local Port: 2249
Remote Port: 3726

Data structure of a socket



Network Byte Order

- All values stored in a `sockaddr_in` must be in network byte order.
 - `sin_port` a port number.
 - `sin_addr` an IP address.

Common Mistake:
Ignoring Network Byte Order



Byte Ordering

```
union {  
    u_int32_t addr; /* 4 bytes address */  
    char c[4];  
} un;  
  
un.addr = 0x8002c25f; /* 128.2.194.95 */  
/* c[0] = ? */
```

- Big Endian

- Sun Solaris, PowerPC, ...

c[0]	c[1]	c[2]	c[3]
128	2	194	95

- Little Endian

- Intel x86, DEC Alpha, ...

95	194	2	128
----	-----	---	-----

- Network byte order = Big Endian

Network Byte Order Functions

'**h**' : host byte order
byte order

'**s**' : short (16bit)

'**n**' : network

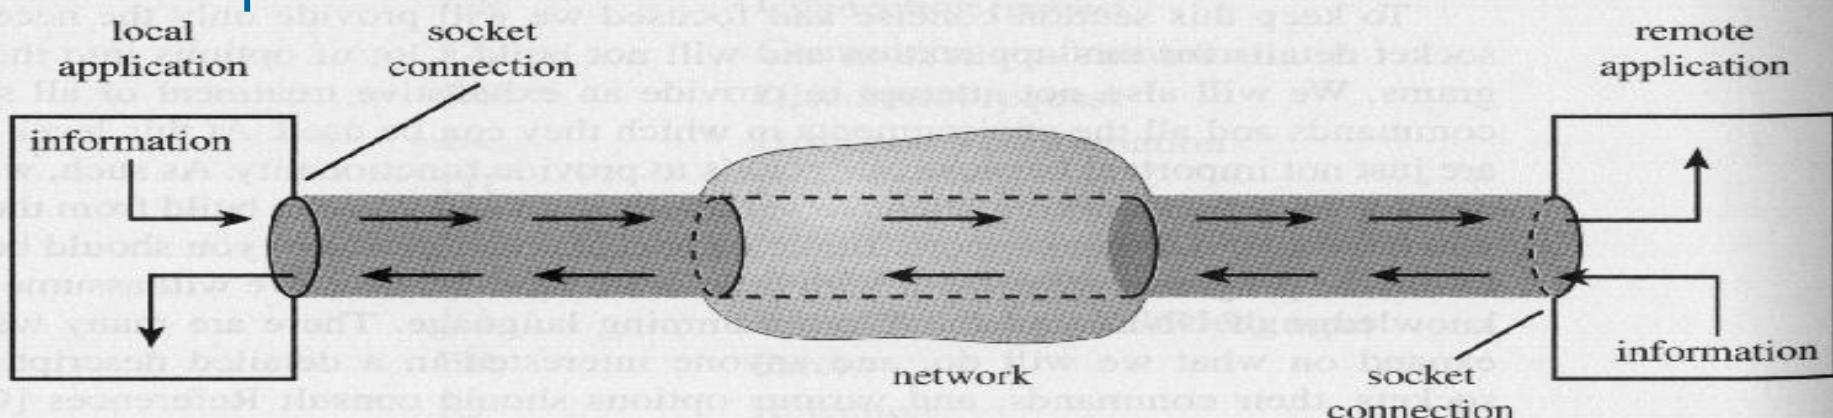
'**l**' : long (32bit)

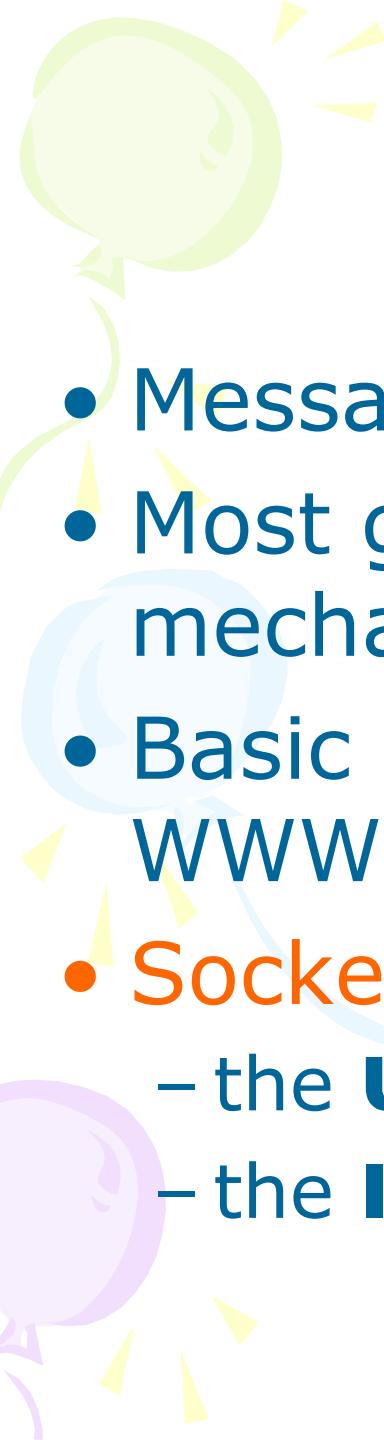
```
uint16_t htons(uint16_t);  
uint16_t ntohs(uint16_t);
```

```
uint32_t htonl(uint32_t);  
uint32_t ntohl(uint32_t);
```

Socket

- Socket is a UNIX construct and is the basis for UNIX network I/O
 - It is available under MS Windows as well, i.e. WinSock 2
- Developed by UC Berkeley: **BSD Socket** (Berkeley Software Distribution)
- Originally used in UCB Unix as an API for TCP (Transport Control Protocol) and UDP (User Datagram Protocol)
- Socket and its related primitives are transport interface to higher layer applications
- Socket is a communication identifier consisting of a local port number and an Internet address





BSD Sockets

- Message passing
- Most general IPC (Inter Process Communication) mechanism
- Basic building block of Internet and WWW
- Sockets have addresses either in
 - the **UNIX domain** (same machine)
 - the **Internet domain**

BSD Sockets

- Three important types of sockets
 - **stream sockets** (TCP Transport Control Protocol) provide reliable delivery of data
 - No data will be lost, replicated or arrive out of sequence
 - **datagram sockets** (UDP User Datagram Protocol) make no such warranties (e.g. ping)
 - **raw sockets:** very low level

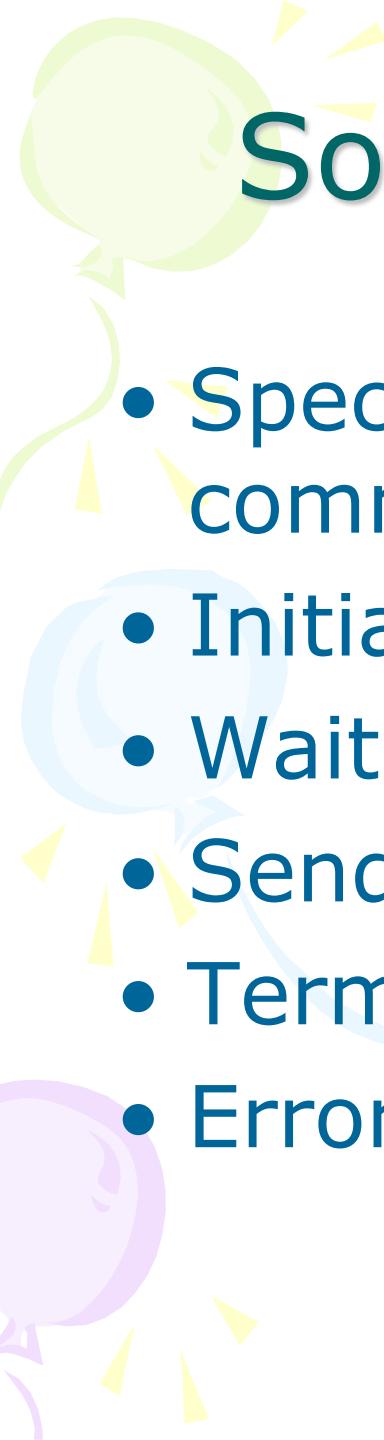
Sockets basics for Linux

Client server applications are often using sockets

- The server provides services to client's programs that connect to the server via socket interface
- Associated with services is a number known as
 - Port
 - 0 -1023 are reserved
 - User code must use numbers < 1023
- Ports combined with the IP address of the computer executing the server program form a unique point of contact (address) for this server.
- For example, the address **114.58.1.6:6072** comprises the computers IP number **114.58.1.6** and the port number **6072** separated by a colon.

/etc/services is a file on most computers used to specify socket addresses

ftp	21/tcp	
Ssh	22/tcp	SSH Remote Login Protocol
Ssh	22/udp	SSH Remote Login Protocol
telnet	23/tcp	
smtp	25/tcp	mail
smtp	25/udp	mail
time	37/tcp	time server
time	37/udp	time server
whois++	63/tcp	nick name
bootps	67/tcp	BOOTP server
bootpc	68/tcp	BOOTP client
gopher	70/tcp	Internet Gopher
gopher	70/udp	
finger	79/tcp	
http	80/tcp	WorldWideWeb HTTP

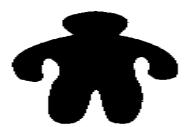


Socket functions needed

- Specify local and remote communication endpoints
- Initiate a connection
- Wait for incoming connection
- Send and receive data
- Terminate a connection gracefully
- Error handling

Socket

- Client/server model
 - Involves two distinct programs running on different machines at different locations
 - They have network connections
 - The server provides services and responds to requests coming in from clients
- Primitives for socket
 - **socket():** create a communication end point
 - **bind():** attach a local address to a socket
 - **listen():** announce number of connections
 - **accept():** waiting for a connection call
 - **connect():** attempt to establish a connection
 - **send()/write():** send data over the connection
 - **recv()/read():** receive data from the connection
 - **close():** close the connection



↑
user makes request
through a client program

Local
machine

runs a
client
program

Network

Remote
machine

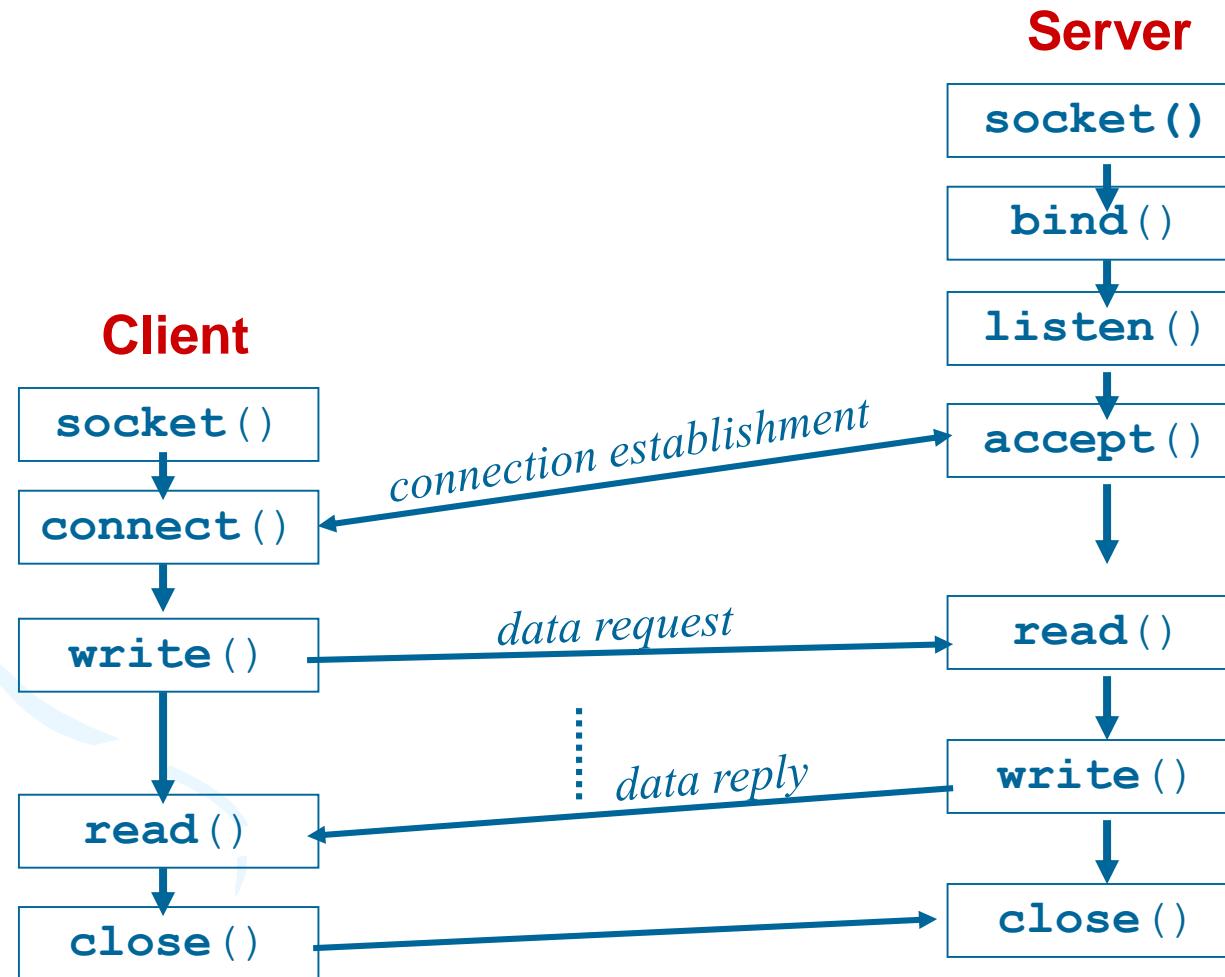
runs a
server
program

↑
server has access
to user files



Auxiliary
storage

Socket Client-Server Interaction



Simple Socket Server

```
#include <stdio.h>
#include <stdlib.h>
#include <netinet/in.h>
#include <sys/socket.h>

int main() {

    int server_socket = socket(AF_INET, SOCK_STREAM, 0);

    struct sockaddr_in server_address, client_address;
    server_address.sin_family = AF_INET;
    server_address.sin_addr.s_addr = INADDR_ANY;
    server_address.sin_port = htons(1234);
    bind(server_socket,(struct sockaddr *)&server_address,sizeof(struct sockaddr_in));

    listen(server_socket, 5);

    socklen_t l = sizeof(struct sockaddr_in); /* client address length */
    int client_socket = accept(server_socket, (struct sockaddr *) &client_address, &l);

    char data[512] = { '\0' };
    int message_length = read(client_socket, data, 512);
    printf("from client: %s\n", data);
    message_length = sprintf(data, "back to ya client");
    message_length = write(client_socket, data, message_length);

    close(client_socket);

    return 0;
}
```

Simple Socket Client

```
#include <stdio.h>
#include <stdlib.h>
#include <netinet/in.h>
#include <sys/socket.h>

int main() {
    int client_socket = socket(AF_INET, SOCK_STREAM, 0);

    struct addrinfo *server_address = NULL;
    struct addrinfo hints;
    memset(&hints, 0, sizeof(struct addrinfo));
    hints.ai_family = AF_INET;
    hints.ai_socktype = SOCK_STREAM;
    getaddrinfo("localhost", "1234", &hints, &server_address);
    connect(client_socket, server_address->ai_addr, server_address->ai_addrlen);

    char data[512] = {'\0'};
    int message_length = sprintf(data, "hello server");
    message_length = write(client_socket, data, message_length);
    message_length = read(client_socket, data, 512);
    printf("from server: %s\n", data);

    close(client_socket);

    return 0;
}
```



Simple Python Socket Server

```
import socket

server_socket = socket.socket(socket.AF_INET, socket.SOCK_STREAM)

server_socket.bind("", 1234)

server_socket.listen(5)

client_socket, address = server_socket.accept()

data = client_socket.recv(512)
print "from client: ", data
data = "back to ya client"
client_socket.send(data)

client_socket.close()
```

Simple Python Socket Client

```
import socket

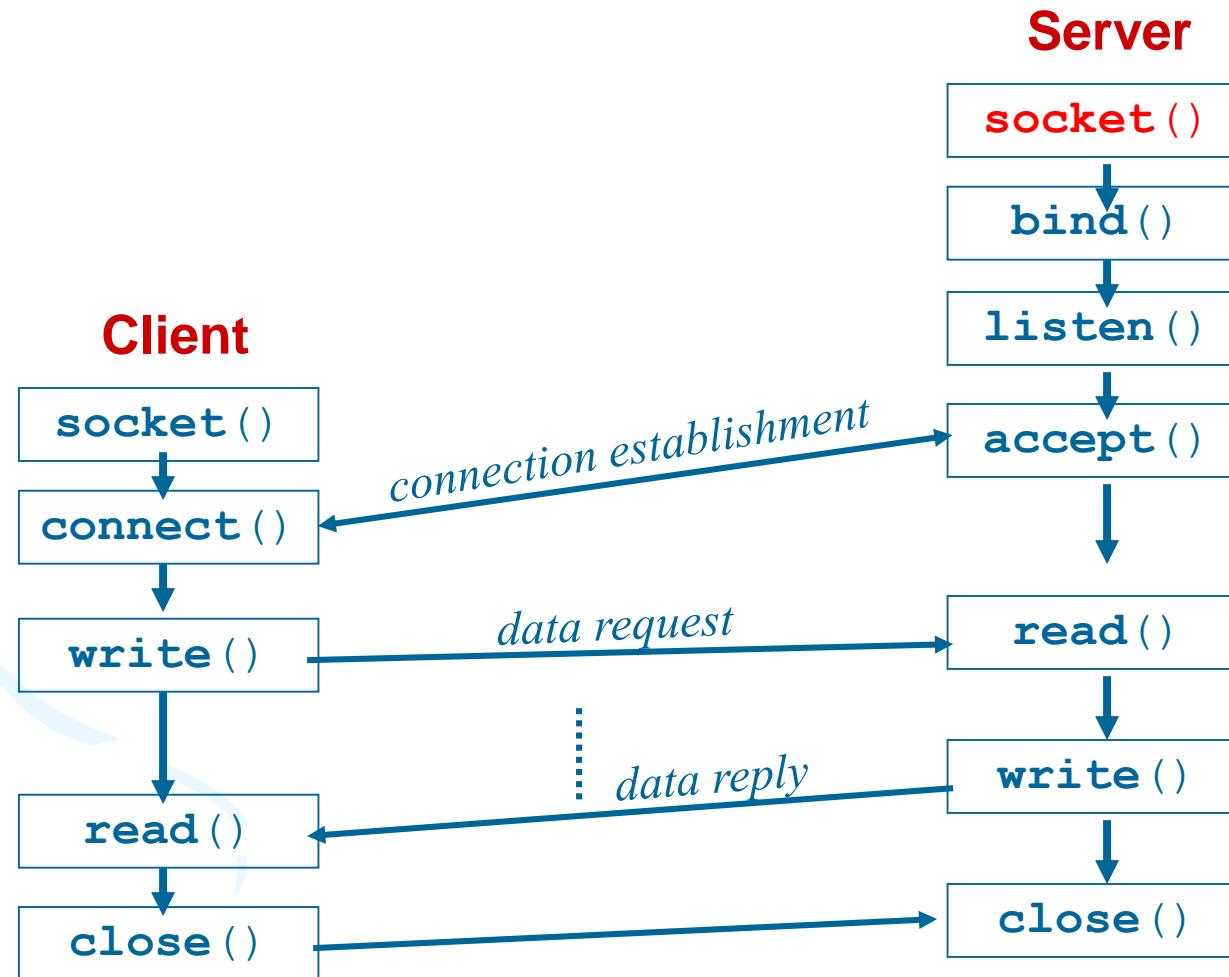
client_socket = socket.socket(socket.AF_INET, socket.SOCK_STREAM)

client_socket.connect(("localhost", 1234))

data = "hello server"
client_socket.send(data)
data = client_socket.recv(512)
print "from server: ", data

client_socket.close()
```

Socket Client-Server Interaction



The socket() function

Prototype:

```
int socket(int family, int type, int protocol)
```

Example – Internet socket:

```
if (( sock = socket(AF_INET,SOCK_STREAM,0))  
== -1)  
{ printf(": error opening socket");  
exit(-1);  
}
```

- Creates a socket:
 - Not associated with a client or server process yet
- Socket type:
 - SOCK_STREAM vs. SOCK_DGRAM

Server code (get socket and initialize struct)

```
main (int argc, char *argv[])
{
    int rc,
        msgsock, sock,
        adrlen,
        cnt;
    /* system call return code */
    /* server/listen socket descriptors */
    /* sockaddr length */
    /* number of bytes I/O */

    struct sockaddr_in
    struct sockaddr_in
    struct sockaddr
    sockname; /* Internet socket name */
    *nptr;    /* ptr to get port number after bind*/
    addr;     /* generic socket name */

    char buf[80];           /* I/O buffer for messages - could be larger*/
    /* You need to lookup in /etc/hosts file. */

    struct hostent *hp, *gethostbyaddr();

/*get socket API*/

    if (( sock = socket(AF_INET, SOCK_STREAM, 0) ) < 0 ) {
        perror("Get server socket");
        exit(1);
    }
    /* Initialize the fields in the Internet socket name structure*/
    sockname.sin_family = AF_INET;          /* Internet address */
    sockname.sin_port = 0;                  /* System will assign port # */
    sockname.sin_addr.s_addr = INADDR_ANY;  /* "Wildcard" */
}
```

```
/*
struct hostent {
    char *h_name;           //official name
    char **h_aliases;       // alias list
    int h_addrtype;         //address type
    int h_length;           // address length
    char **h_addr_list;     // address list
#define h_addr h_addr_list[0]
    /*backward compatibility*/
};
```

The structures *sockaddr* , *sockaddr_in*

>> **sockaddr**

Contain address information associated
with a socket (generic structure)

>> **sockaddr_in** applies to Internet protocols
(contains *IP address and port*)

They have the same length in bytes

Unix domain sockets

```
struct sockaddr_un {  
    short sun_family; //AF_UNIX  
    char sun_path (108); //path  
};
```

```
struct sockaddr {  
    u_short sa_family;  
    u_short sa_data[14];  
}
```

```
struct sockaddr_in {  
    short sin_family;  
    u_short sin_port;  
    struct in_addr sin_addr;  
    char sin_zero[8]; // unused  
}
```

Using sockaddr_in structure

Example I:

```
sock_addr.sin_addr.s_addr = htonl(INADDR_ANY);  
sock_addr.sin_port = htons(0);  
sock_addr.sin_family = AF_INET;
```

Example II:

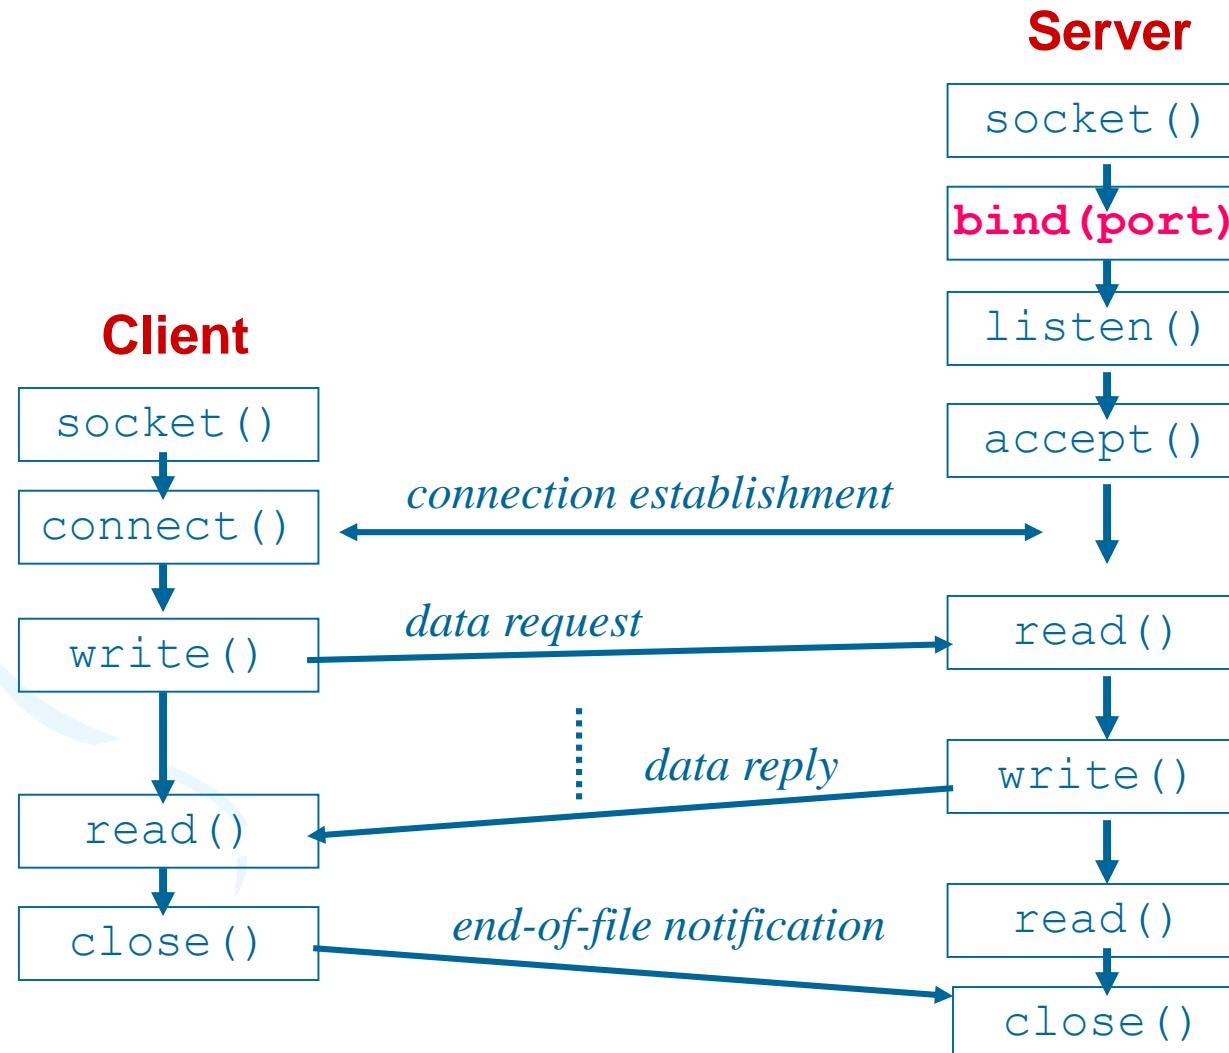
```
sock_addr.sin_addr.s_addr = inet_addr("127.59.69.7");  
sock_addr.sin_port = htons(5115);  
sock_addr.sin_family = AF_INET
```

Example 3: we need more:

```
struct sockaddr_in sa;  
struct sockaddr_in caller;  
struct sockaddr_in *nptr; /*ptr to get port number*/  
struct sockaddr_in addr; /*generic socket address*/
```

- **htonl(), htons()** handle byte ordering differences between computer architectures and network protocols
 - *redundant in Sun's Sparc*
 - *needed in x86 architectures*
- **INADDR_ANY** is used for accepting connections on any local interface if the system is multi-homed
- We assign a zero port number when filling the **sockaddr_in** data structure, if we want a free port to be allocated to a socket during the bind() procedure

Socket Client-Server Interaction



The bind() function

```
int bind(int sockfd, struct sockaddr *myaddr, int addrlen)
```

Example:

```
if ( bind(sock, (struct sockaddr *) &sock_addr, sizeof(sock_addr)) < 0)
    { printf(": error binding socket to local address");
      exit(-1);
    }
```

This call binds a socket to an IP address and port.
It is not associated with a client or server process yet

1. **sock** is a **file descriptor** for the socket
2. **name** is a **pointer to a structure** of type **sockaddr**

Server code – bind() and get port number

```
if (bind(sock, &sockname, sizeof(sockname)) < 0) {
    close(sock);                                /* always close socket */
    perror("network server bind");
    exit(2);
}

/*Get the port number assigned to the Internet socket and
print it for use as arg to client*/

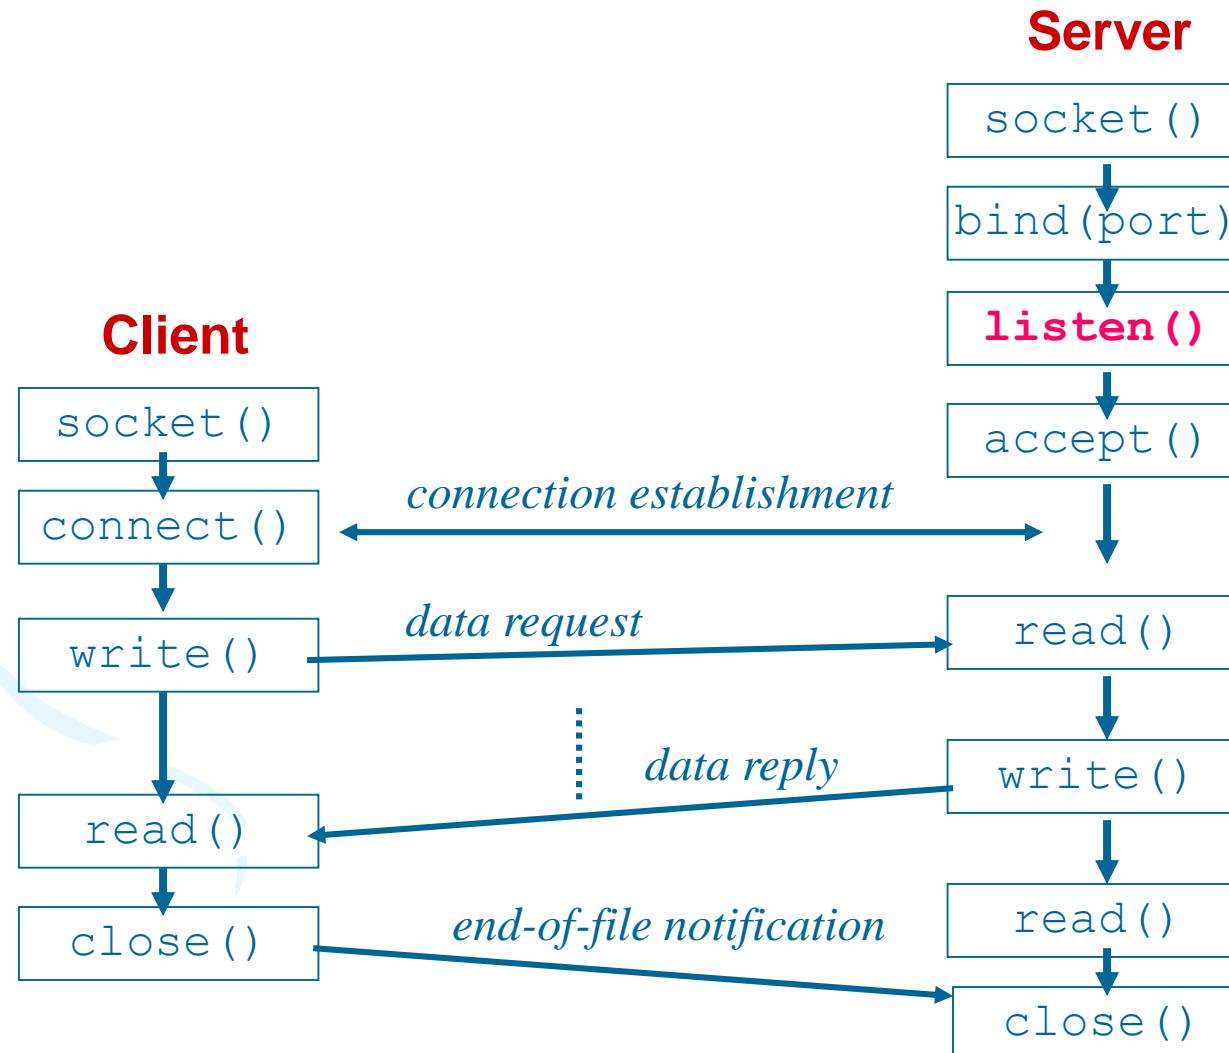
adrflen = sizeof(addr);
if ( ( rc = getsockname( sock, &addr, &adrflen ) ) < 0 )
{
    perror("network server getsockname");
    close (sock);
    exit(3);
}

/*tell the client the socket number.
The pointer nptr points to the generic address structure. */

nptr = (struct sockaddr_in *) &addr;      /* port number */
                                         /* now you can print port number is for client to use */
```

socket port is in nptr->sin_port >>> see next slide

Socket Client-Server Interaction



Function listen()

```
/*get the port assigned to the Internet socket and print it */

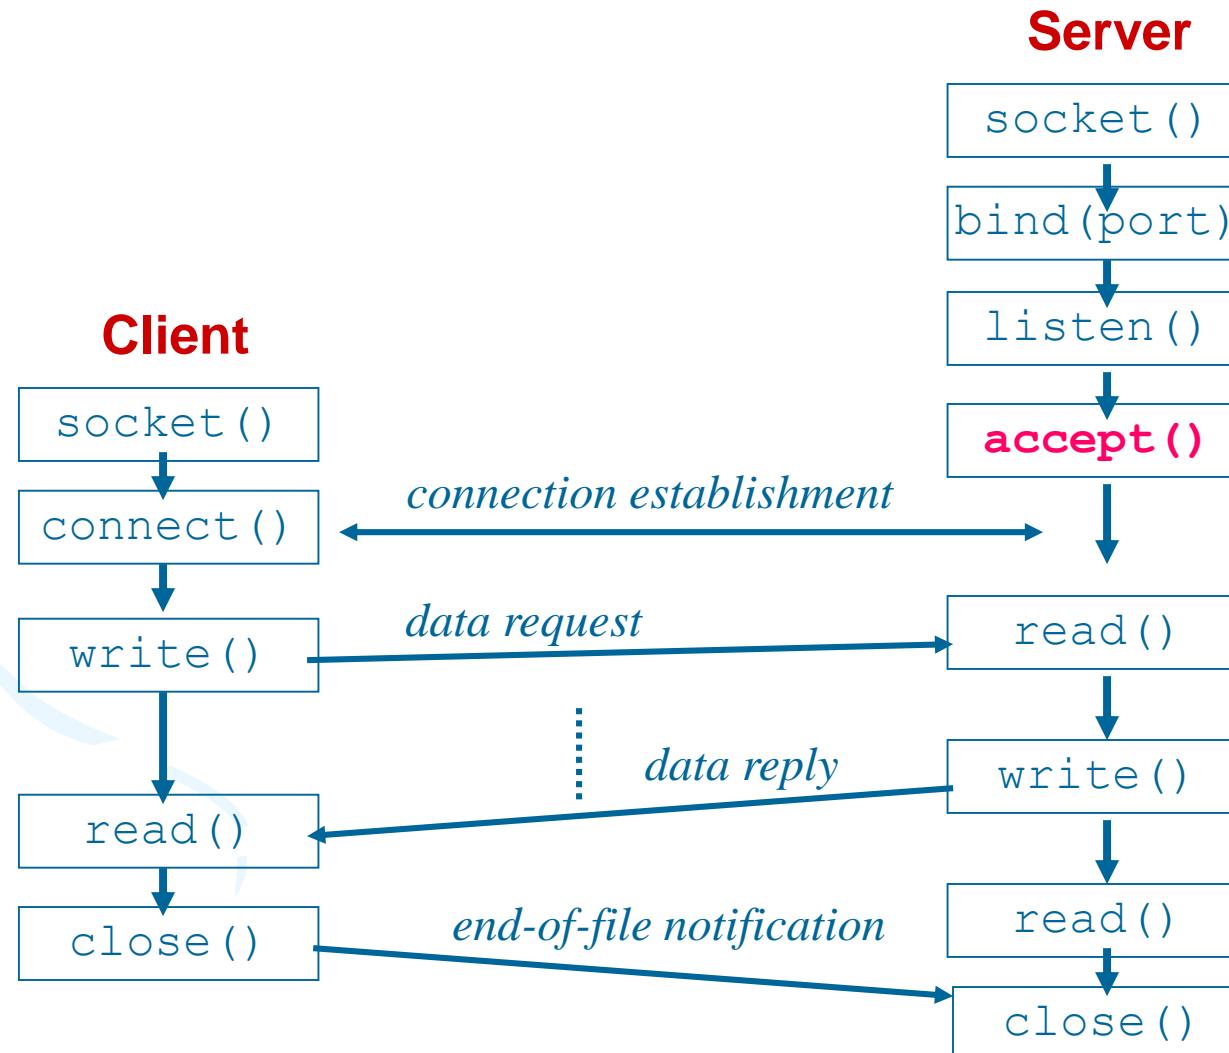
length = sizeof(sa);      /*length of entire structure!!*/
/*check if function worked */
if ((rc=getsockname(s, (struct sockaddr_in *)&sa, &length) )< 0) {
    close(s);
    ERROR("getsockname");
}
nptr = (struct sockaddr_in *) &sa;      /*get port number */
printf("\n\Server has port number: %d\n", ntohs(nptr->sin_port));

/* listen for connections on a socket */

listen (s, 5);
```

/*listen establishes the socket as a passive endpoint of a connection.
It does not suspend the process execution; No messages can be sent through
this socket, however, incoming messages can be received
sock is a file descriptor associated with the socket*/

Socket Client-Server Interaction



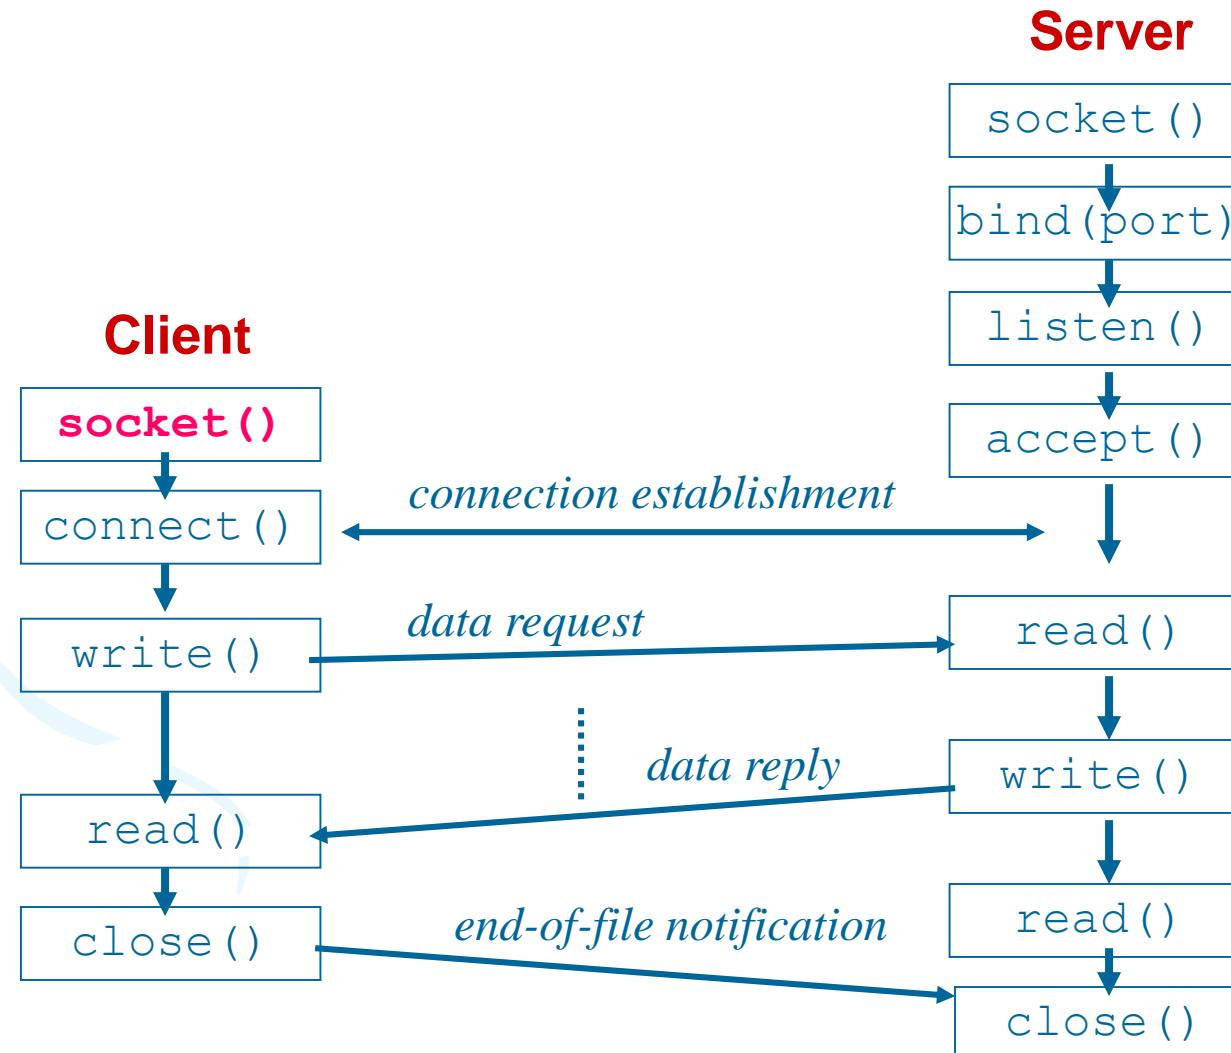


Function accept()

```
/* Use accept to wait for calls to the socket. Accept returns
*a new socket which is connected to the caller. msgsock will be a
 *temporary (non-reusable) socket different from s
*msgsock is returned if the call is successful */
```

```
length = sizeof(caller);
if ((msgsock = accept(s, (struct sockaddr *)&caller,&length)) < 0) {
    printf("\nNetwork server accept failure %d\n", errno);
    perror("network server");
    close(s);
    exit(5);
....}
```

Socket Client-Server Interaction



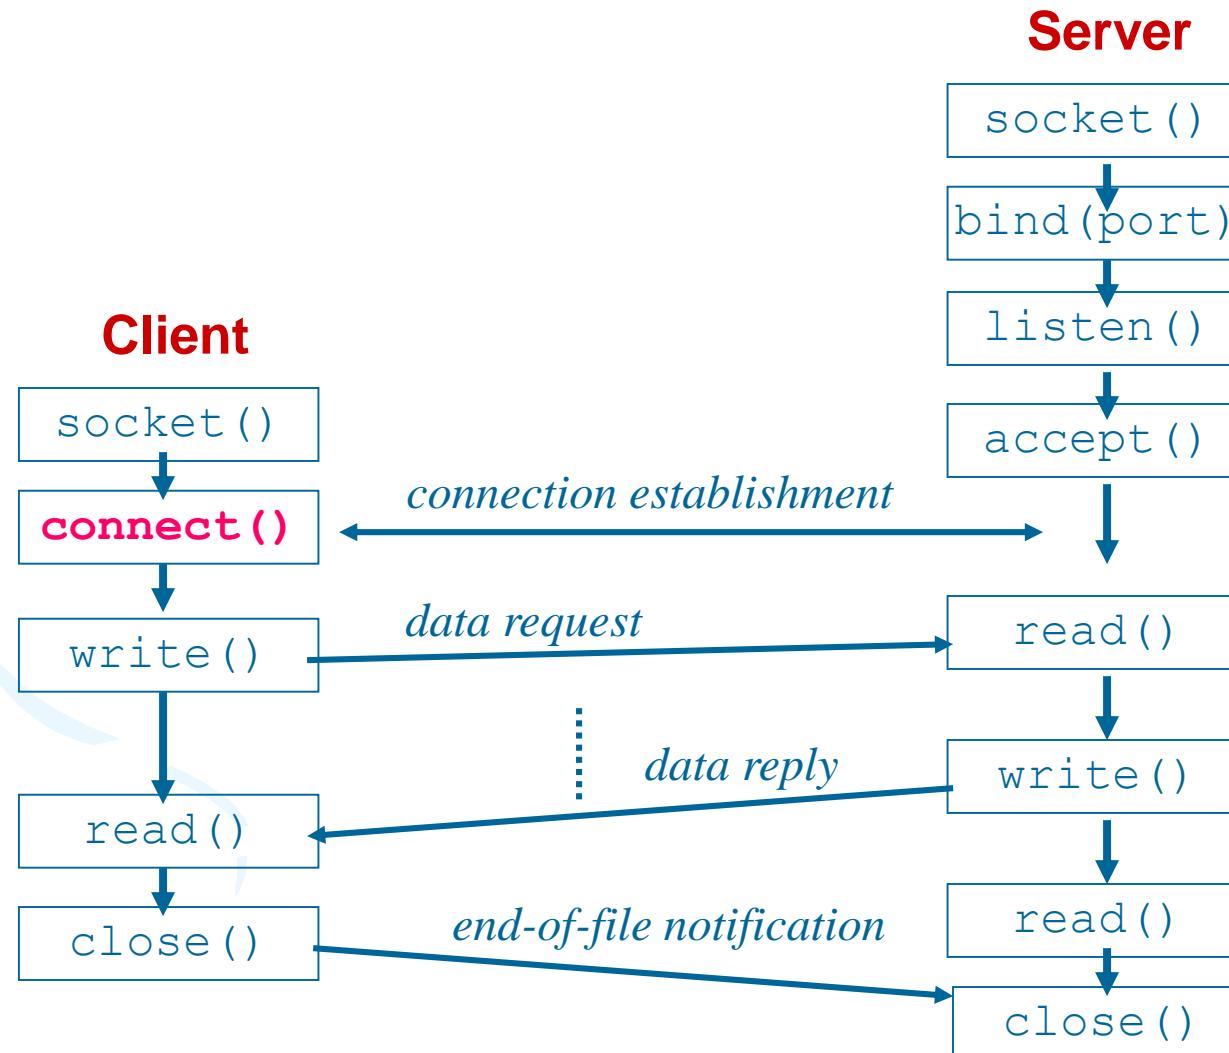
Client programming

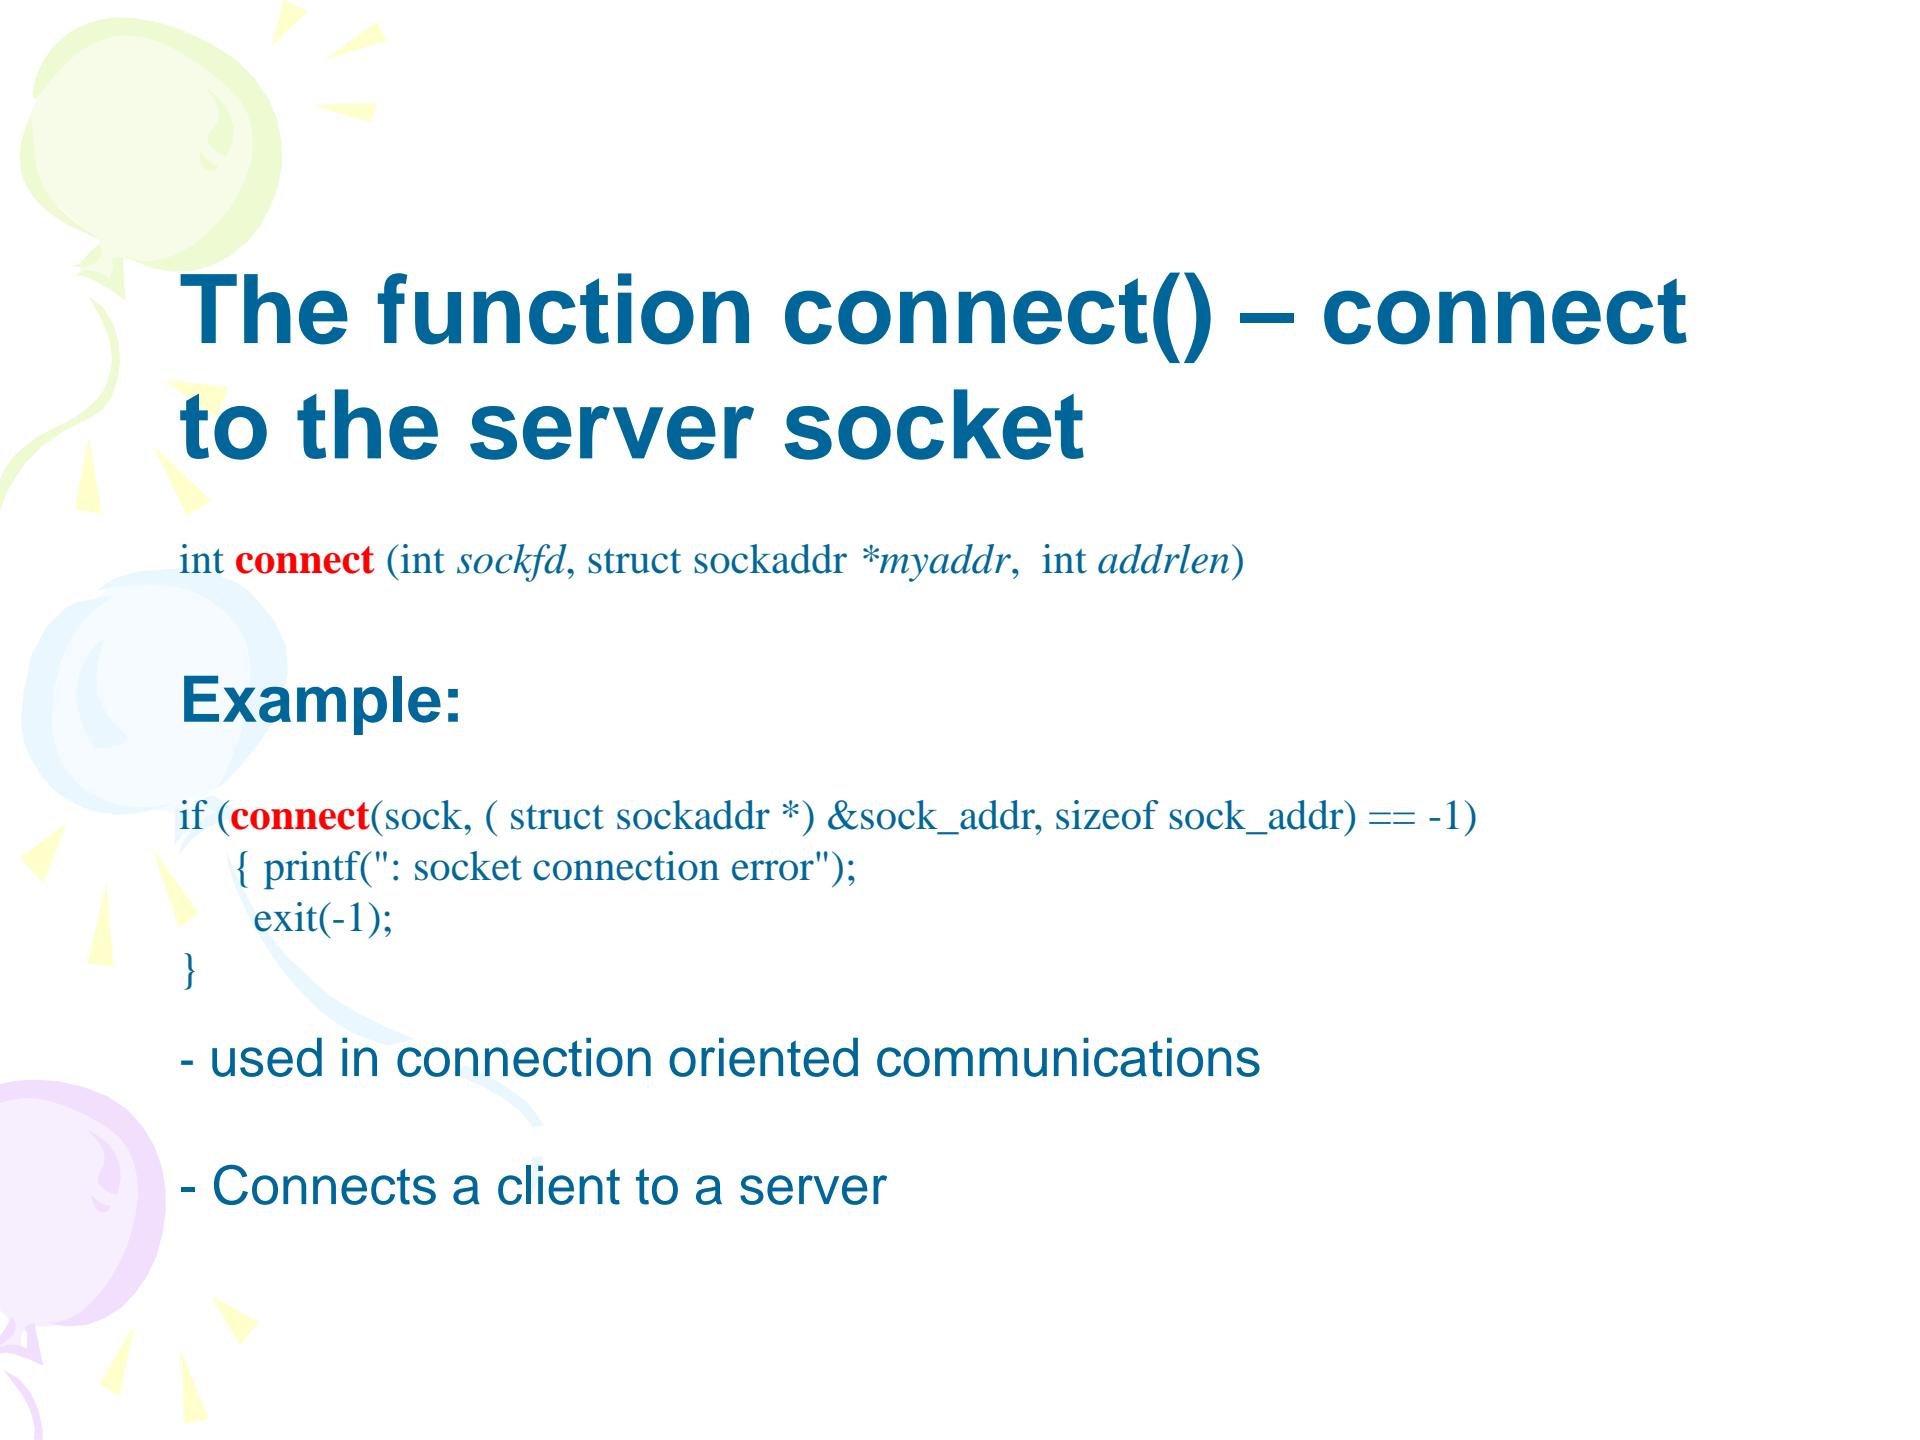
```
/* Expected parameters on command line:  
  
    argv[0] -- name of executable (client)  
    argv[1] -- the host name to which connection is desired (use localhost)  
    argv[2] -- the port number to be used by the client: the  
              value is the port number assigned to the  
              server by the server's host system. */  
  
{  
    int sock,           /* socket descriptor */  
    val,                /* scratch variable */  
    cnt;                /* number of bytes I/O */  
  
    struct sockaddr_in client; /* Internet socket name (addr) */  
    struct sockaddr_in *nptr; /* pointer to get port number */  
  
    char buf[80];      /* I/O buffer, for the message - could be larger */  
  
/* For lookup in /etc/hosts file. */  
    struct hostent *hp, *gethostbyaddr();
```

Client code to establish the socket

```
{  
    int sock,                      /* socket descriptor */  
        val,                         /* scratch variable */  
        cnt;                         /* number of bytes I/O messages*/  
  
    struct sockaddr_in client;     /* Internet socket name (addr) */  
    struct sockaddr_in *nptr;      /* pointer to get port number */  
  
    char buf[80];                  /* I/O buffer, kind of small , could be larger*/  
  
    /* For lookup in /etc/hosts file. */  
    struct hostent *hp, *gethostbyaddr();  
  
    if (( sock = socket(AF_INET, SOCK_STREAM, 0)) < 0 ) {  
        perror("network client");  
        exit(2);  
    }  
    /* Convert user-supplied port number to integer. */  
  
    client.sin_port = htons( atoi(argv[2]) );           /* Server port number */  
    client.sin_family = AF_INET;                        /* Internet domain */
```

Socket Client-Server Interaction





The function `connect()` – connect to the server socket

```
int connect (int sockfd, struct sockaddr *myaddr, int addrlen)
```

Example:

```
if (connect(sock, ( struct sockaddr *) &sock_addr, sizeof sock_addr) == -1)
    { printf(": socket connection error");
      exit(-1);
    }
```

- used in connection oriented communications
- Connects a client to a server

Client can get the host name

```
hp = gethostbyname (argv[1]);  
  
/* It is good to know to whom we are connected to */  
if ( hp == NULL ) {  
    perror("network client");  
    close (sock);  
    exit(3);  
}  
else {  
    printf("\tThe official host name is %s\n", hp -> h_name);  
    printf("\tThe first host address is %lx\n",  
          ntohl ( * (int * ) hp -> h_addr_list[0] ) );  
    printf("\tAlias names for the host are:\n");  
  
    while ( *hp -> h_aliases )           //get all aliases  
        printf( "\t\t%s\n", *hp -> h_aliases++ );  
}  
  
bcopy ( hp -> h_addr_list[0], &client.sin_addr.s_addr,  
          hp -> h_length ); /* save the first name */
```

Connection between server and client

```
/* Client: Establish socket connection with (remote) server. */

if ( ( connect ( sock, &client, sizeof(client) ) ) < 0 ) {
    printf("network client %s connect failed %d\n", argv[0], errno);
    perror("network client on connect");
    close (sock);
    exit(4);
}
```

listen()

```
/* SERVER: Set up "infinite loop" to listen for clients. Since the
   structure "sockname" is bound to the listen socket, the
   socket structure name and socket length parameter
   values are omitted from the accept call. The bound values
   are used. */

while (1) {
    if ( ( msgsock = accept (sock, 0, 0 ) ) < 0 ) {
        printf("Network server accept failed %d\n", errno);
        perror("network server accept");
        close (sock);
        exit(5);
    }
    else { /*processing of the call comes here) /*}
}
```

The read() and write() functions

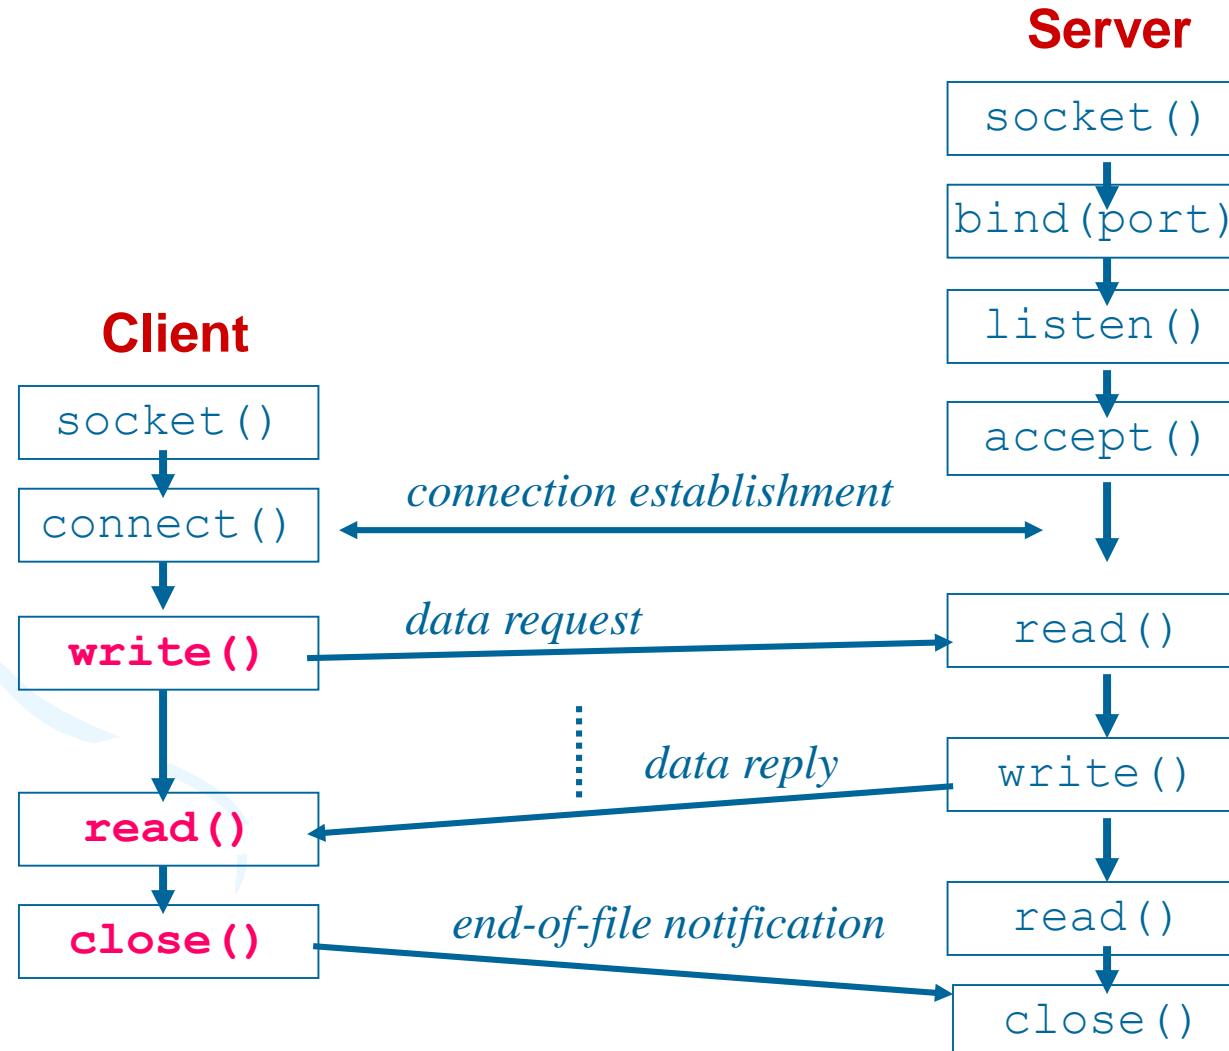
```
int read(int sockfd, char *buf, unsigned int nbytes)
int write(int sockfd, char *buf, unsigned int nbytes)
```

Examples:

```
if (write(sock, message, size) < 0 )
{ printf(": error writing to remote host");
  exit(-1);
}
bytes_received = read(sock, reply, sizeof(reply));
```

- used in connection oriented communications
- used for requesting and transmitting data

Socket Client-Server Interaction



Client established the communication

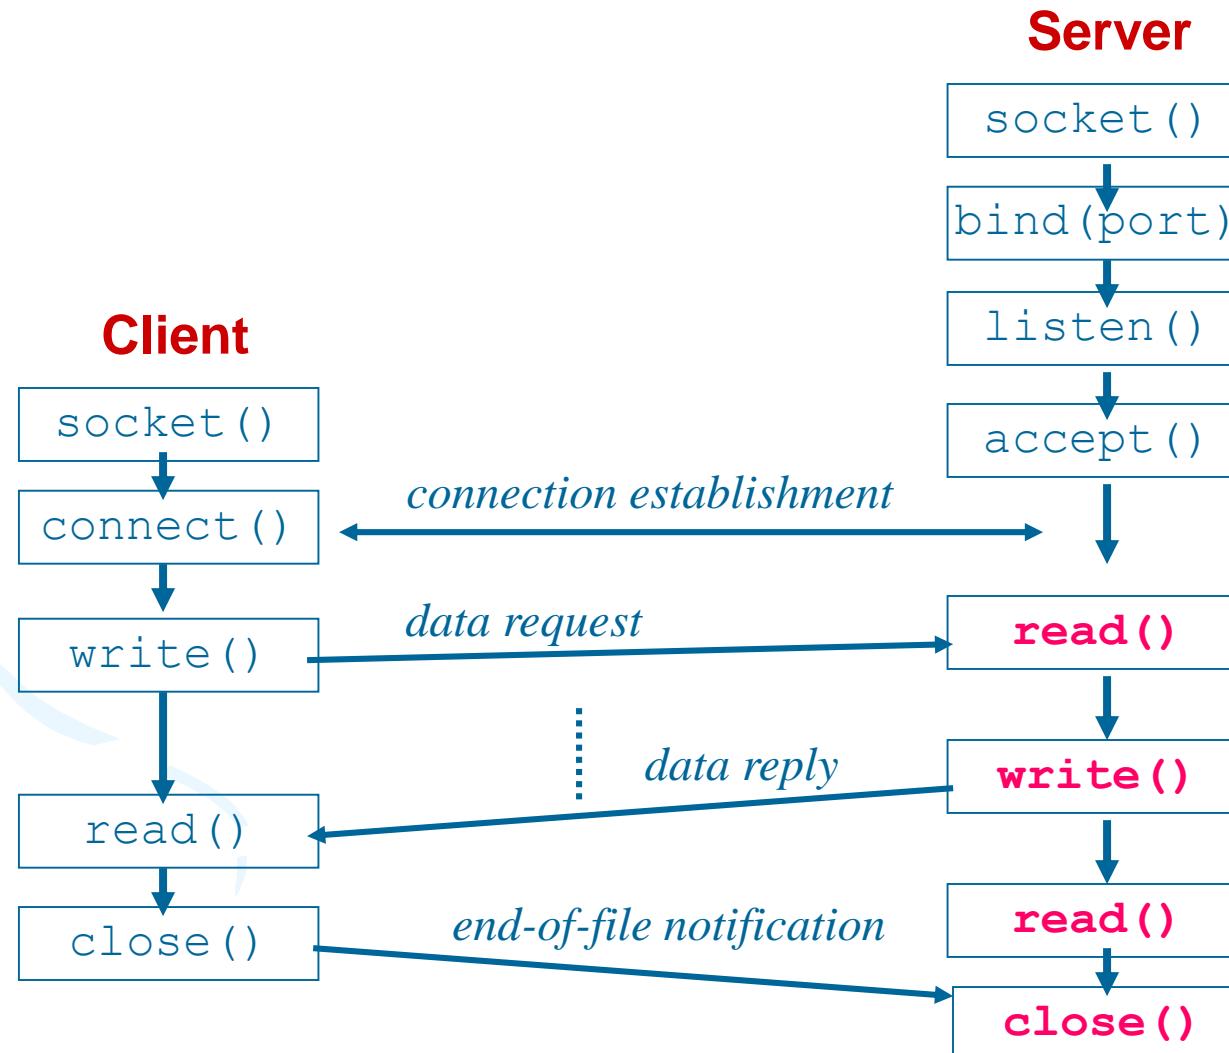
```
if ( ( connect ( sock, &client, sizeof (client) ) ) < 0 ) {
    perror ("network client on connect");
    close (sock);
    exit(4);
}
/* Exchange data with server. Clear buffer bytes first. */

bzero ( buf, sizeof( buf) ); /* zero buffer, BSD. */
strcpy ( buf, "Hello server, please respond to client!" );
write ( sock, buf, sizeof(buf) );

/* Now read message sent back by server. */

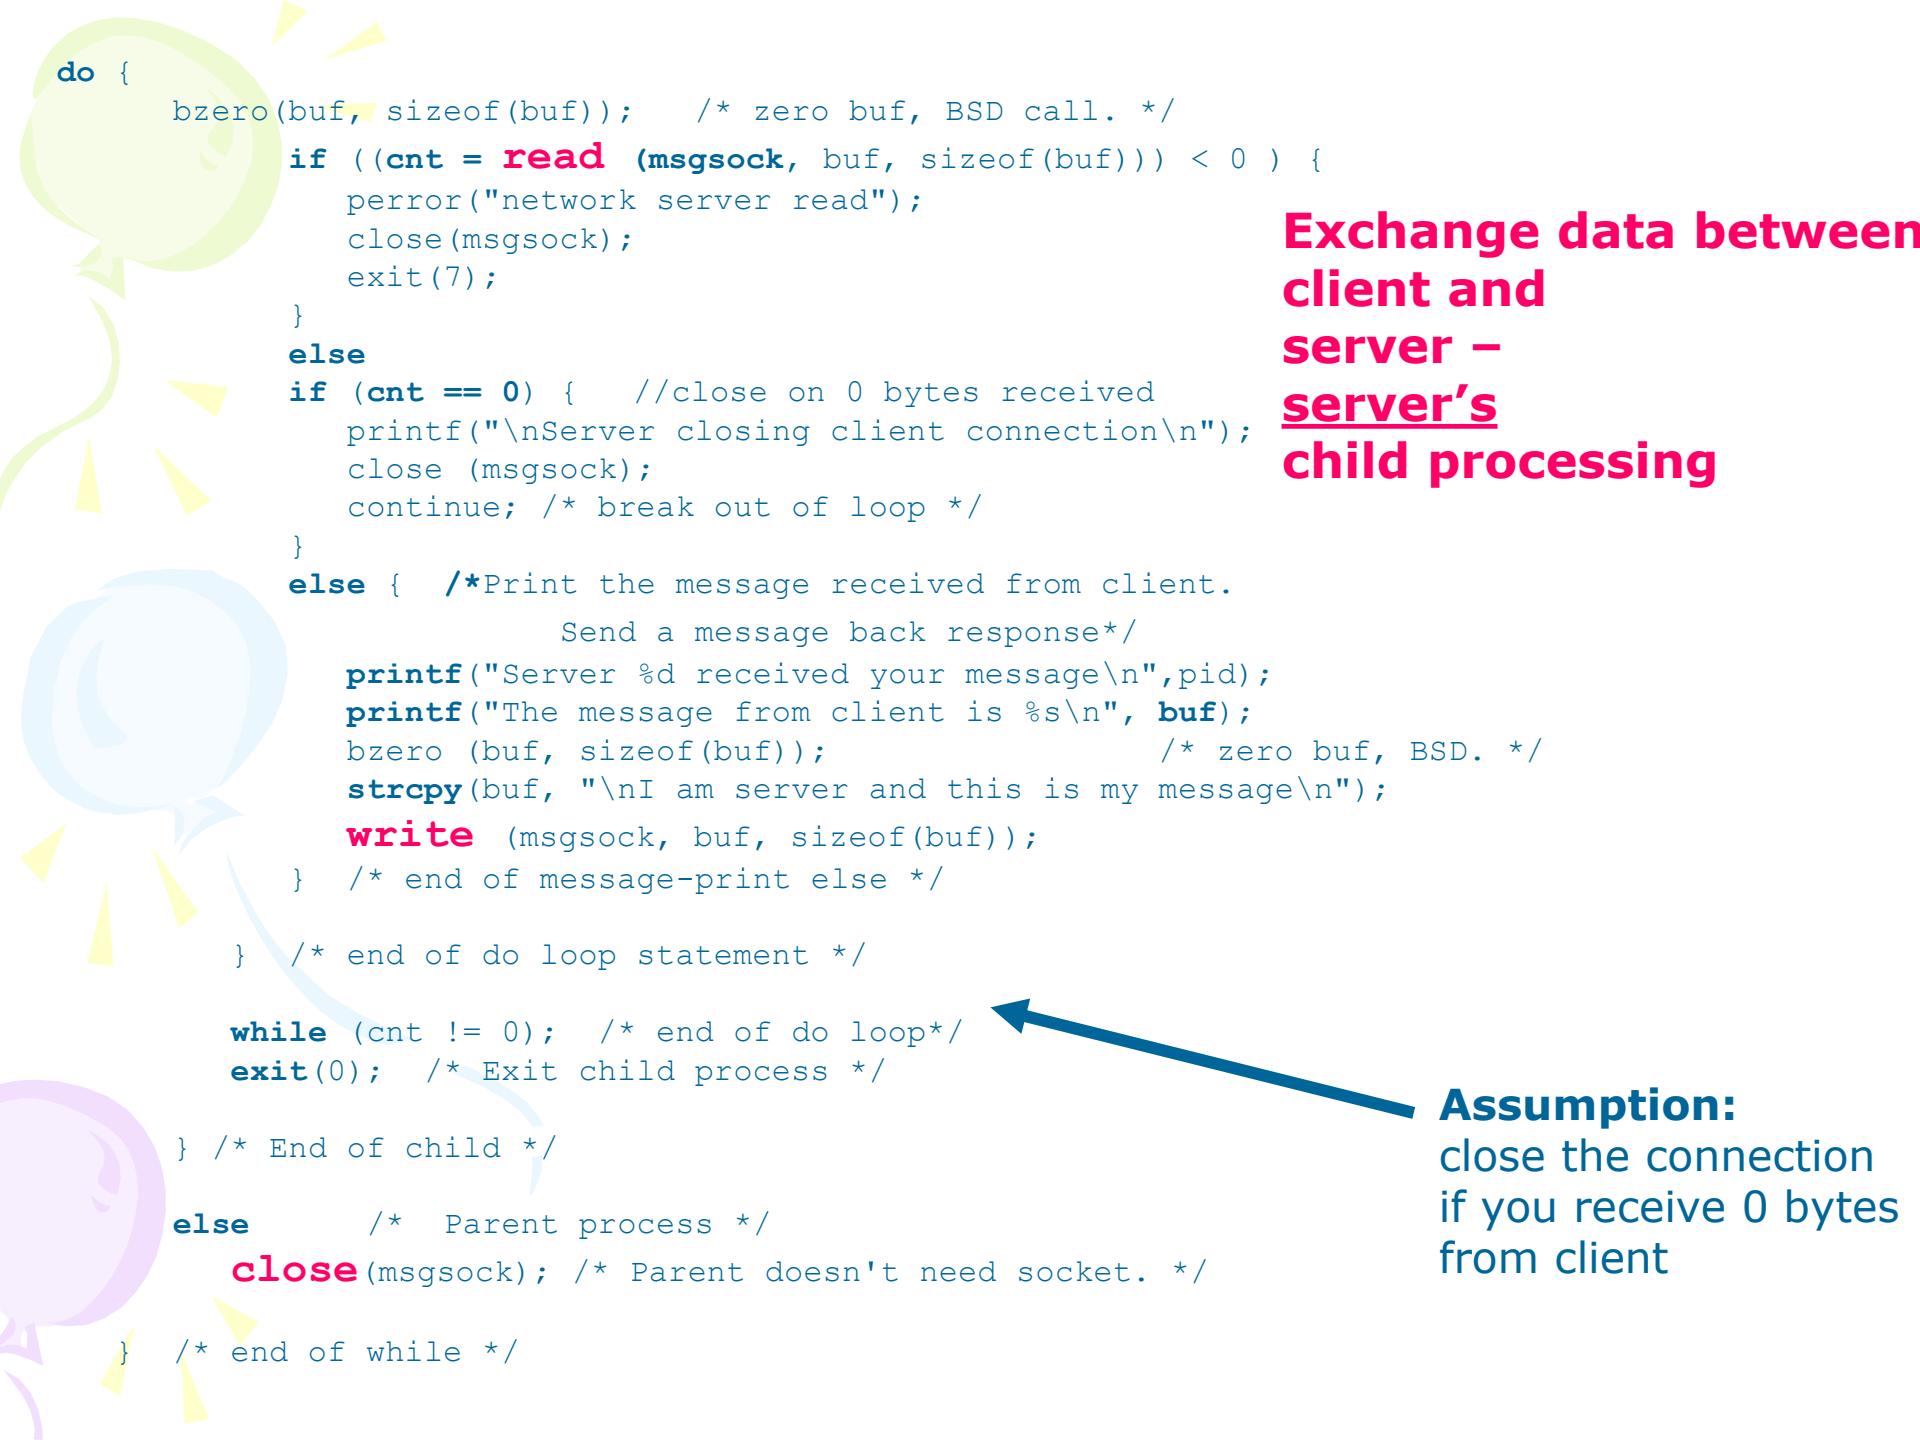
if ( ( cnt = read (sock, buf, sizeof(buf) ) ) < 0 ) {
    perror("network client on read");
    close(sock);
    exit(5);
}
else
    printf("I am client and I have received this message %s\n", buf);
    /* Send a message with 0 bytes to end the conversation. */
bzero ( buf, sizeof( buf) ); /* zero buffer, BSD. */
write ( sock, buf, 0 );
close (sock);
exit(0);
```

Socket Client-Server Interaction



Server communication

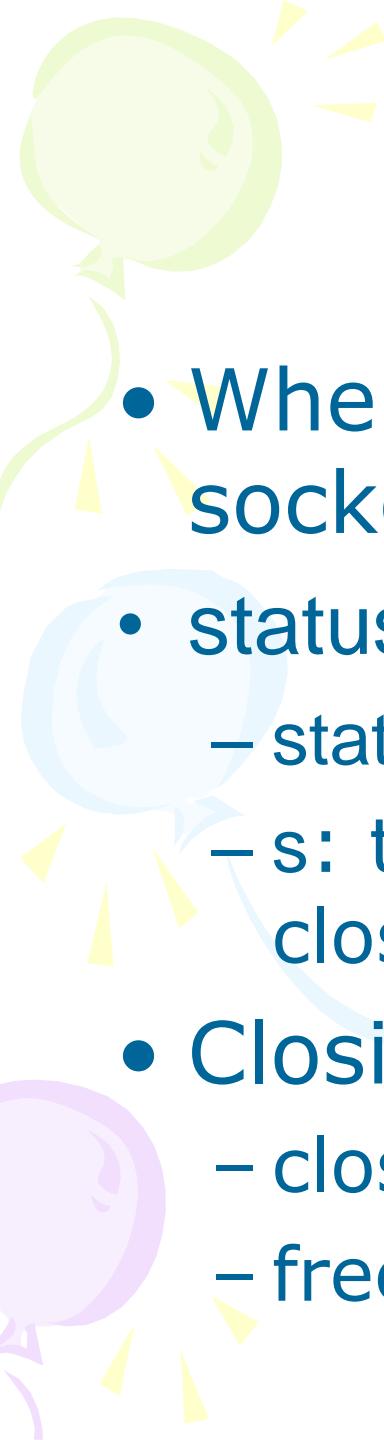
```
/* Server communication includes:  
Set up "infinite loop" to listen for clients (see slide 21).  
The bound values are used. */  
  
while (1) {  
    if ( ( msgsock = accept (sock, 0, 0) ) < 0 ) {  
        perror("network server accept");  
        close (sock);  
        exit(5);  
    }  
  
    /* Fork child process to handle client service request */  
  
    if ( ( fork() ) == 0 ) {      /* Child process to handle communication*/  
        int pid;  
        pid = getpid();          /* PID of child process */  
        close (sock);            /* Do not need listen socket in child. */  
        /* Find out who the client is. */  
        /* */  
        if ((rc = getpeername( msgsock, &addr, &adrlen )) < 0) {  
            perror("network server getpeername");  
            close(msgsock);  
            exit(6);  
        }  
        printf("\n\tnetwork server %d:", pid);  
        printf(" client socket from host %s\t has port number %d\n",inet_ntoa ( nptr -> sin_addr ),  
              nptr -> sin_port );
```



Exchange data between client and server – server's child processing

```
do {  
    bzero(buf, sizeof(buf)); /* zero buf, BSD call. */  
    if ((cnt = read (msgsock, buf, sizeof(buf))) < 0 ) {  
        perror("network server read");  
        close(msgsock);  
        exit(7);  
    }  
    else  
    if (cnt == 0) { //close on 0 bytes received  
        printf("\nServer closing client connection\n");  
        close (msgsock);  
        continue; /* break out of loop */  
    }  
    else { /*Print the message received from client.  
            Send a message back response*/  
        printf("Server %d received your message\n",pid);  
        printf("The message from client is %s\n", buf);  
        bzero (buf, sizeof(buf)); /* zero buf, BSD. */  
        strcpy(buf, "\nI am server and this is my message\n");  
        write (msgsock, buf, sizeof(buf));  
    } /* end of message-print else */  
}  
/* end of do loop statement */  
  
while (cnt != 0); /* end of do loop*/  
exit(0); /* Exit child process */  
  
} /* End of child */  
  
else /* Parent process */  
close(msgsock); /* Parent doesn't need socket. */  
} /* end of while */
```

Assumption:
close the connection
if you receive 0 bytes
from client



The close() function

- When finished using a socket, the socket should be closed:
- `status = close(s);`
 - `status`: 0 if successful, -1 if error
 - `s`: the file descriptor (socket being closed)
- Closing a socket
 - closes a connection (for SOCK_STREAM)
 - frees up the port used by the socket

Socket Server: Python code with equivalent C code

```
import socket
#include <stdio.h>
#include <stdlib.h>
#include <netinet/in.h>
#include <sys/socket.h>
int main() {

    server_socket = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
    int server_socket = socket(AF_INET, SOCK_STREAM, 0);

    server_socket.bind("", 1234)
    struct sockaddr_in server_address, client_address;
    server_address.sin_family = AF_INET;
    server_address.sin_addr.s_addr = INADDR_ANY;
    server_address.sin_port = htons(1234);
    bind(server_socket,(struct sockaddr *)&server_address,sizeof(struct sockaddr_in));

    server_socket.listen(5)
    listen(server_socket, 5);

    client_socket, address = server_socket.accept()
    socklen_t l = sizeof(struct sockaddr_in); /* client address length */
    int client_socket = accept(server_socket, (struct sockaddr *) &client_address, &l);

    data = client_socket.recv(512)
    print "from client: ", data
    data = "back to ya client"
    client_socket.send(data)
    char data[512] = { '\0' };
    int message_length = read(client_socket, data, 512);
    printf("from client: %s\n", data);
    message_length = sprintf(data, "back to ya client");
    message_length = write(client_socket, data, message_length);

    client_socket.close()
    close(client_socket);
    return 0;
}
```

Socket Client: **Python code** with equivalent **C code**

```
import socket
#include <stdio.h>
#include <stdlib.h>
#include <netinet/in.h>
#include <sys/socket.h>
int main() {

    client_socket = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
    int client_socket = socket(AF_INET, SOCK_STREAM, 0);

    client_socket.connect(( "localhost", 1234))
    struct addrinfo *server_address = NULL;
    struct addrinfo hints;
    memset(&hints, 0, sizeof(struct addrinfo));
    hints.ai_family = AF_INET;
    hints.ai_socktype = SOCK_STREAM;
    getaddrinfo( "localhost", "1234", &hints, &server_address);
    connect(client_socket, server_address->ai_addr, server_address->ai_addrlen);

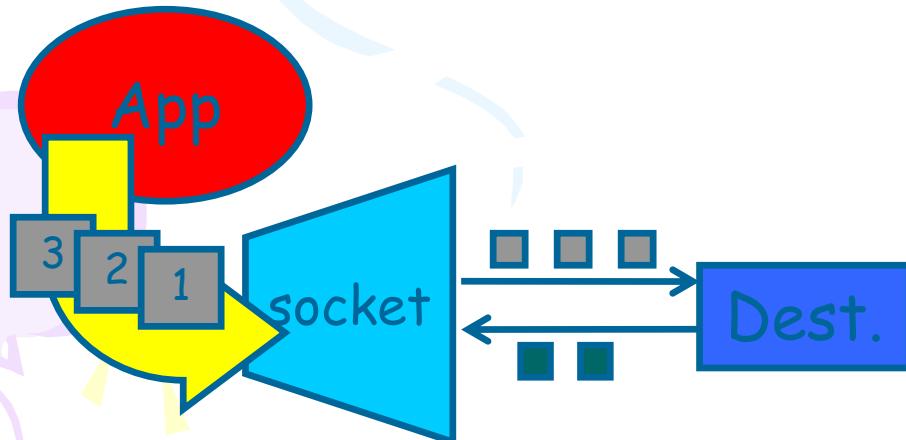
    data = "hello server"
    client_socket.send(data)
    data = client_socket.recv(512)
    print "from server: " , data
    char data[512] = {'\0'};
    int message_length = sprintf(data, "hello server");
    message_length = write(client_socket, data, message_length);
    message_length = read(client_socket, data, 512);
    printf("from server: %s\n", data);

    client_socket.close()
    close(client_socket);
    return 0;
}
```

Two essential types of sockets

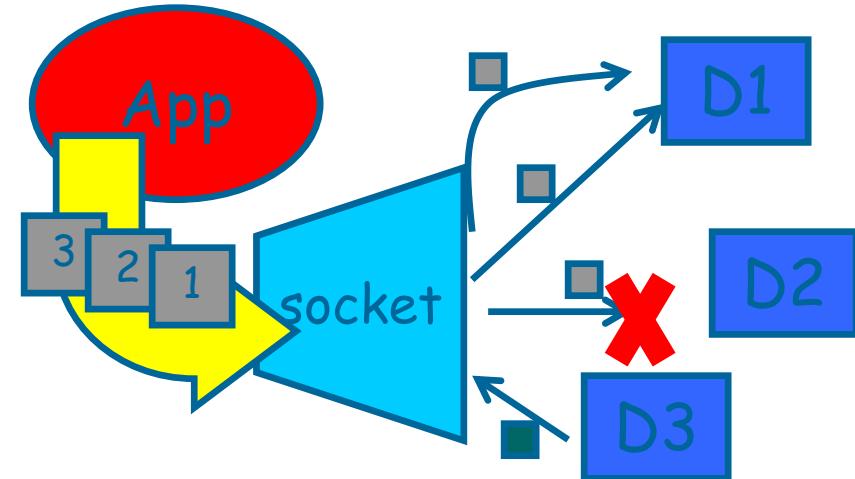
stream socket

- a.k.a. TCP
(Transport Control Protocol)
- reliable delivery
- in-order guaranteed
- connection-oriented
- bidirectional

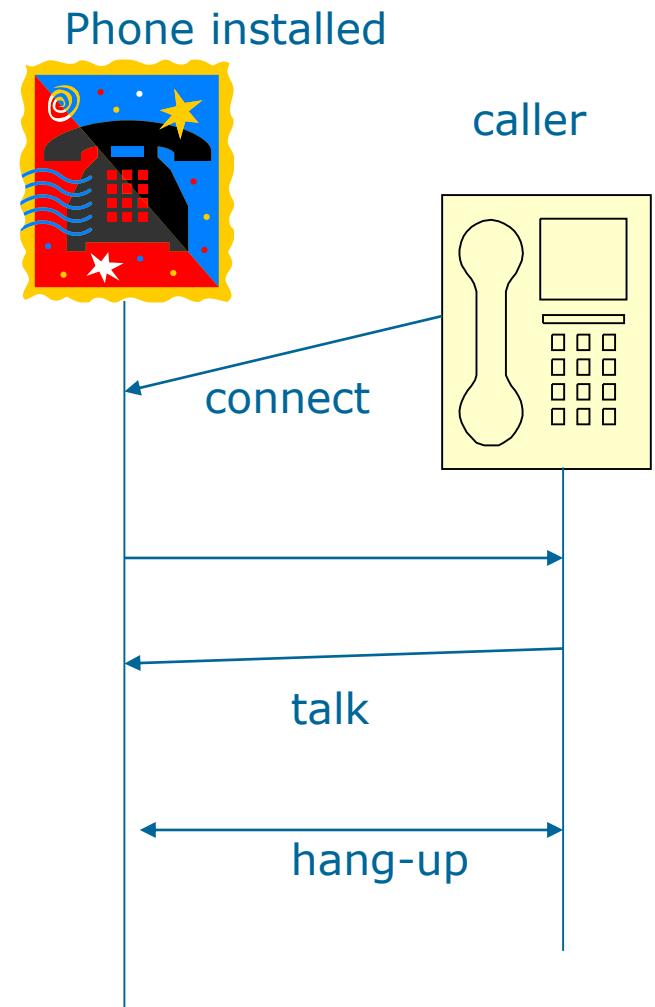
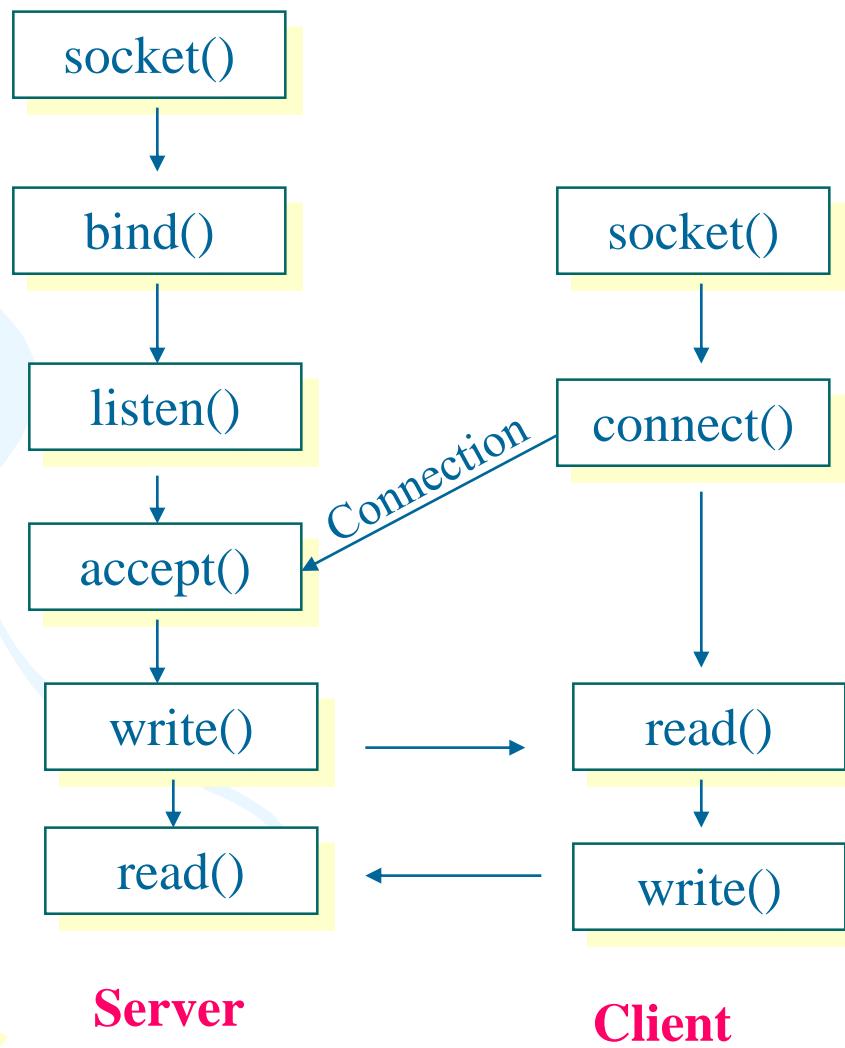


datagram socket

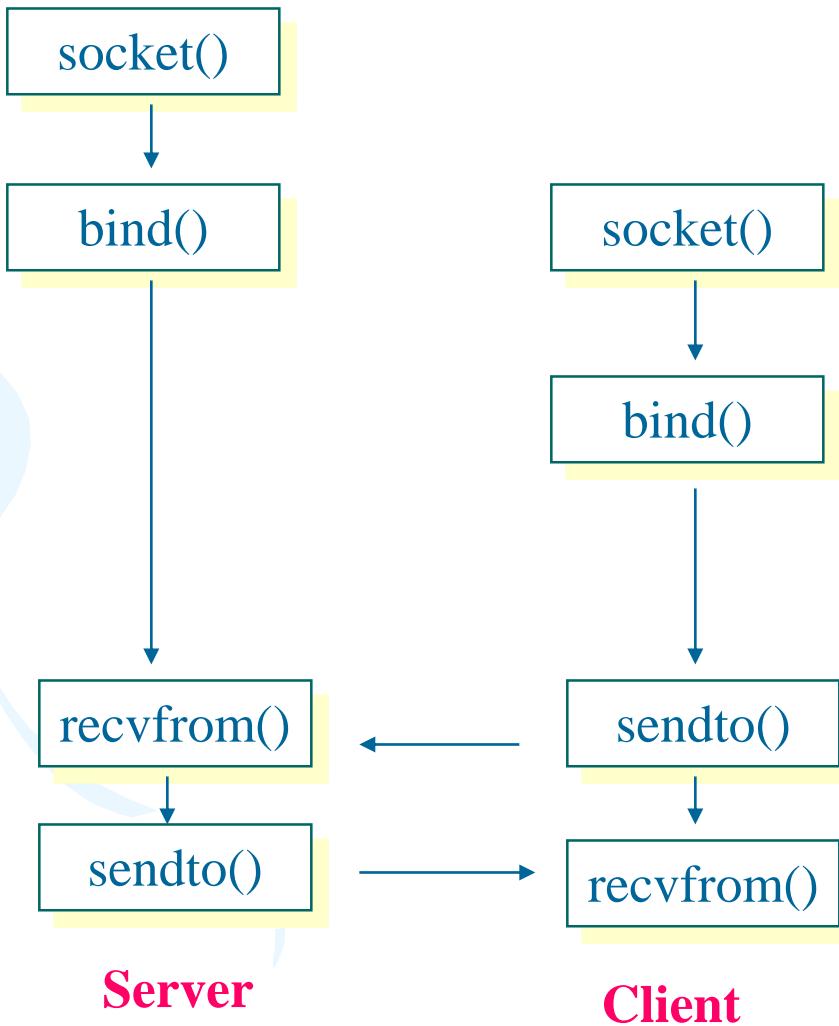
- a.k.a. UDP (e.g. ping)
(User Datagram Protocol)
- unreliable delivery
- no order guarantees
- no notion of “connection” where app indicates a destination for each packet
- can send or receive



Connection oriented protocol = stream socket



Connectionless Protocol = datagram socket



Sockets

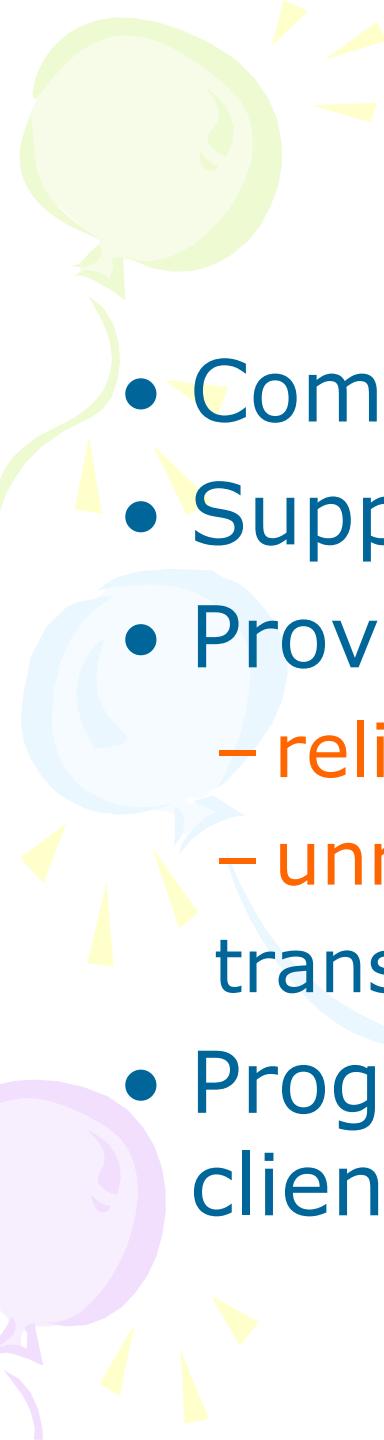
- **Comments:**

Streams vs Datagrams: If a sender makes 5 separate calls to write, each for 512 bytes and the receiver asks for 2560 bytes then in the case of a

- **stream:** the receiver gets all the requested bytes in the correct order
- **datagram:** only 512 bytes is returned - four other calls will be needed to retrieve all 2560 bytes

Datagrams are less frequently used than streams.

Datagrams are used primarily for services that send small amounts of data on short time connections, or send only a little amount of data at a time.



Summary Sockets are:

- Communication abstractions
- Support the TCP/IP protocol stack
- Provide access to both
 - reliable (TCP Transport Control Protocol) and
 - unreliable (UDP User Datagram Protocol)transport services
- Programming tools to implement client-server applications

Sample Exam Question

List eight differences between stream and datagram sockets.

- **Stream Socket**
- uses TCP
- bidirectional
- (synchronous)
- stream data (large)
- connection orientated
- (persistent connection)
- (need to close connection)
- in order guaranteed
- (so message boundaries not ..)
- reliable delivery
- (acknowledge receipt)
- more complex
- application: HTTP, FTP

vs

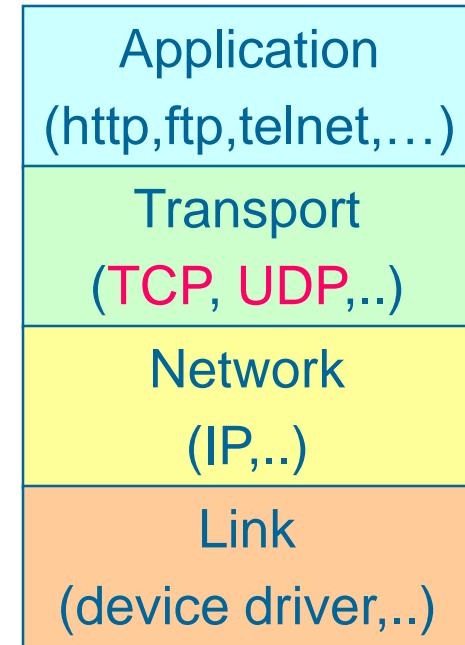
Datagram Socket

- uses UDP
- one way only – send or receive
- (asynchronous)
- single packet (small)
- not connection orientated
- (single-use connection)
- (no connection to close)
- no order of data guaranteed
- (message boundaries preserved)
- unreliable delivery
- (do not acknowledge receipt)
- simpler – less overhead
- application: ping

Networking Basics

- Applications Layer
 - Standard apps
 - HTTP
 - FTP
 - Telnet
 - User apps
- Transport Layer
 - TCP (Transport Control Protocol)
 - UDP (User Datagram Protocol)
 - Programming Interface:
 - Sockets
- Network Layer
 - IP
- Link Layer
 - Device drivers

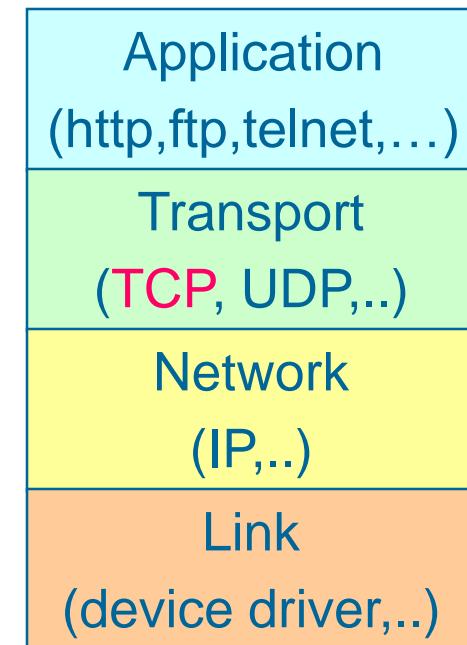
TCP/IP Stack:



Networking Basics

- TCP (Transport Control Protocol) is a connection-oriented protocol that provides a reliable flow of data between two computers.
- Example applications:
 - HTTP
 - FTP
 - Telnet

TCP/IP Stack:



Networking Basics

- UDP (User Datagram Protocol) is a protocol that sends independent packets of data, called *datagrams*, from one computer to another with no guarantees about arrival.
- Example applications:
 - Clock server
 - Ping

TCP/IP Stack:

